

Financing the Space-Based Energy Transition: Endogenous Risk Mitigation, Spatial Competition, and Sustainable Mechanism Design

Shiming Wu

Abstract

The staggering computational requirements of artificial intelligence are creating an unprecedented terrestrial energy and thermal bottleneck. While mega-scale space solar energy remains a long-term goal, technology firms and sovereign actors are currently engaged in a pre-emptive orbital competition for AI data computation centers. Because near-Earth orbits are a finite, common-pool resource in a global game where international treaties are unenforceable, this paper develops a continuous-time stochastic differential game to analyze this astropolitical market failure caused by externality. I first prove mathematically why sovereign states will rationally defect from centralized space treaties, necessitating private, self-enforcing financial contracts. I then design a mechanism—the orbital-sustainability bond—and provide step-by-step Hamilton-Jacobi-Bellman derivations for the firm’s valuation under endogenous, state-dependent jump-to-default risk. By expanding the framework into structural corporate finance, I solve for staged compound real options and Power Purchase Agreement (PPA) derivative collars that mathematically cure the agency friction between founders and institutional bondholders. The optimally calibrated mechanism generates a double dividend: it lowers the Phase 1 entry threshold to $P = 3.46$, while heavily delaying secondary competitors and inducing a spatial pivot to deep space. Ultimately, this framework provides a highly actionable corporate finance architecture that halts the value-destroying pre-emptive race and yields a 21.7% recovery of deadweight loss.

1 Introduction

The advancement of frontier technology is fundamentally bounded by physical energy capacity. The computational demands of training and operating large computing models have created a severe terrestrial energy bottleneck. A single next-generation chip consumes approximately 64.8 kilowatt-hours daily, and with technology firms planning thousands of new facilities, digital infrastructure is projected to account for a significant double-digit percentage of all United States electricity usage by the end of the decade. This demand shock is colliding with a rigid terrestrial grid, prompting sovereign nations and multinational corporations to direct capital toward space-based solar power and extraterrestrial energy generation, which promises uninterrupted, baseload renewable power.

However, this transition introduces a fundamental problem of industrial organization, mechanism design, and astropolitics. The optimal near-Earth orbits and strategic planetary locations represent finite common-pool resources. A critical aspect of these resources is that congestion is strictly endogenous to total spatial occupancy. A monopolist scaling infrastructure to meet consumer demand inherently increases its own risk of triggering an orbital cascade. Furthermore, when a second firm enters the same orbital sector, the difficulties of inter-firm operational coordination cause the risk of a terminal Kessler Syndrome collapse to spike exponentially. Kessler Syndrome refers to a theoretical scenario where the density of objects in low Earth orbit is high enough that collisions between objects cause an irreversible kinetic cascade, rendering specific orbital ranges inaccessible for generations.

A natural regulatory response to common-pool resource degradation is a centralized Pigouvian tax or a cap-and-trade system. However, extraterrestrial domains present a unique jurisdictional challenge. Governed loosely by the 1967 Outer Space Treaty, no single sovereign possesses the authority to enforce mandates or levy taxes in orbital space. This enforcement vacuum, coupled with the escalating physical risks, motivates the core economic research questions of this paper. First, in a continuous-time stochastic game where collapse risk depends non-linearly on spatial occupancy, why does standard market competition inevitably lead to a Pareto-inefficient pre-emptive race? Second, how can sustainable finance mechanisms, acting as private commitment devices, internalize these spatial externalities without state intervention? Third, does the introduction of sustainable finance always yield optimal results, or can poorly calibrated green contracts paradoxically accelerate the tragedy of the commons? Fourth, under what precise empirical conditions does an asymmetric greenium generate a double dividend that aligns private technological expansion with the preservation of the astropolitical commons? Fifth, how does the introduction of deep space deployment alternatives transform the pre-emptive race into a problem of spatial arbitrage? Finally, can

financial contracts mathematically induce private capital investment into reusable rockets and debris capture technology to actively suppress the physical failure rate?

The ultimate objective of this spatial transition extends beyond merely beaming power to terrestrial grids. As artificial intelligence models scale exponentially, they face absolute physical limits on Earth regarding thermal dissipation and land utilization. Consequently, this paper posits that the optimal long-term equilibrium involves migrating the data centers themselves into the extraterrestrial commons. By co-locating artificial intelligence computation centers directly with space-based solar arrays, firms can utilize the near absolute zero temperature of the space vacuum for costless thermal management. This paradigm shift permanently alleviates terrestrial global warming vectors and guarantees long-term energy safety by entirely removing industrial computational loads from the fragile Earth biosphere. While deploying computation to deep space currently exhibits higher capital expenditures than utilizing terrestrial deserts or nuclear facilities, the extraterrestrial paradigm represents an infinite growth horizon that mathematically surpasses the rigid physical boundaries of the Earth.

Answering these questions and financing this transition requires a sophisticated integration of game theory and structural corporate finance. Deploying \$14 billion of unhedged capital to low Earth orbit represents a massive bottleneck, as technology companies and venture capitalists are unwilling to hold unproven infrastructure risk of this magnitude on single balance sheets. To bridge this gap, I extend the continuous-time framework to incorporate compound real options for staged deployment and Special Purpose Vehicle (SPV) capital structures. By utilizing Modigliani-Miller theories alongside Merton-style jump-to-default credit models, I demonstrate how debt tranching and Green Bond issuance can distribute risk. Furthermore, I mathematically prove that when a technology giant executes a Power Purchase Agreement (PPA) with an embedded price floor, it acts identically to an American put option. This derivative structure synthesizes capital, effectively acting as an off-balance-sheet equity injection that prevents SPV insolvency and guarantees the debt tranche. Crucially, this financial architecture neutralizes the asset substitution problem, aligning the incentives of subordinated equity holders with senior ESG bondholders and making the \$14 billion infrastructure transition bankable.

This preservation aims to restore Kaldor-Hicks efficiency, a criterion utilized to evaluate whether a change in resource allocation improves aggregate social welfare. Under standard Kaldor-Hicks theory, a transition is considered efficient if the aggregate gains to the winners are large enough that they could, in theory, compensate the losers, resulting in a net social improvement. While traditional macroeconomic models often leave this compensation as a theoretical footnote, this paper proposes a structural financial mechanism that transforms

this potential into a tangible transfer of wealth. By internalizing the spatial externality, the mechanism recovers immense deadweight loss and contractually captures a defined portion of that surplus—projected in this calibration at \$410 million—for a global equity fund aimed at terrestrial climate adaptation.

Regarding methodology, I develop a continuous-time real options game incorporating Poisson jump-diffusion processes to model state-dependent spatial collapse. The parameters of the stochastic process are explicitly calibrated using maximum likelihood estimation on wholesale electricity data from the EIA and PJM Interconnection. Furthermore, I integrate empirical power purchase agreement pricing from LevelTen Energy to differentiate intermittent solar pricing from baseload pricing. Endogenous orbital failure probabilities are parameterized using European Space Agency conjunction assessments. I solve the Hamilton-Jacobi-Bellman equations to identify the value of the project across three distinct regimes: monopoly scarcity, duopoly abundance, and terminal collapse. The model is subsequently extended to incorporate mutually exclusive spatial options driven by a secondary geometric Brownian motion representing future extraterrestrial city demand, alongside a static optimization problem where firms dynamically select capital allocations for debris mitigation. Furthermore, to address the reality that space infrastructure requires external funding, I construct a structural model of venture capital contracting to establish endogenous equity dilution thresholds. Finally, I utilize a two-factor partial differential equation framework to integrate the stochastic decline of rocket launch costs alongside rising energy demand.

The quantitative results of this paper demonstrate that an unconstrained market yields a symmetric pre-emptive race where firms invest prematurely. Specifically, while a theoretical unthreatened monopolist would wait for an energy price threshold of \$8.37/MWh, the threat of competition forces initial entry at a depressed threshold of \$5.33/MWh, severely cannibalizing expected surplus. I find that the orbital-sustainability bond completely resolves this through asymmetric finance. Protected by an empirically grounded 50 basis-point greenium and a 1.474x baseload revenue premium, the first-mover accelerates clean energy deployment, dropping the entry threshold to \$2.98/MWh. Simultaneously, the heavy contractual penalty imposes a Pigouvian delay on the secondary entrant, driving the follower’s threshold from \$16.78 out to \$21.81/MWh.

The extended framework reveals that this heavy terrestrial penalty induces a spatial pivot, mathematically redirecting the secondary firm to deploy deep space artificial intelligence centers to capture the long-term growth option of extraterrestrial urbanization. Furthermore, by making physical risk endogenous to capital expenditure, I prove the mechanism forces the firm to invest an optimal \$34.0 million equivalent into reusable systems, which drops the physical failure rate and sustainably lowers the long-term entry barrier back to

\$12.10/MWh. Through rigorous mathematical derivations of the capital structure, I demonstrate that sovereign wealth co-investments can substantially reduce required founder dilution, preventing premature funding failures. Ultimately, stochastic simulations prove that this comprehensive mechanism secures the orbital commons and yields a 21.7% total welfare recovery ratio.

The remainder of this paper proceeds as follows. Section 2 reviews the literature. Section 3 provides the institutional background. Section 4 constructs the mathematical environment and the baseline pre-emptive race. Section 5 introduces the orbital-sustainability bond mechanism. Section 6 explores spatial arbitrage, while Section 7 quantifies terrestrial macroeconomic spillovers. Section 8 models endogenous risk mitigation. Section 9 details the financial architecture, encompassing staged deployment, SPVs, and PPA collars. Section 10 solves the two-factor boundary problem. Section 11 outlines the empirical calibration. Section 12 discusses numerical results, and Section 13 concludes.

2 Literature Review

This paper bridges four distinct strands of economic literature: continuous-time real options and game theory, the economics of space and orbital debris, environmental macroeconomics, and structural sustainable finance.

First, this research builds fundamentally upon the real options literature concerning strategic investment under uncertainty. Foundational work by [Dixit and Pindyck \(1994\)](#) established the mathematical apparatus for evaluating irreversible investments. This framework was subsequently extended into game-theoretic oligopoly settings by [Grenadier \(2002\)](#), [Huisman \(2013\)](#), and [Smit and Trigeorgis \(2012\)](#) to analyze how the threat of competitive entry accelerates investment, leading to rent dissipation and pre-emptive investment races. A critical gap in this existing literature is that models of pre-emptive races typically focus strictly on market share competition and price cannibalization, assuming the underlying physical asset remains perfectly intact regardless of competitor entry. This literature generally ignores the physical degradation of shared common-pool resources. My paper contributes to this field by integrating a permanent, state-dependent jump-risk collapse directly into the continuous-time pre-emption framework, explicitly deriving how endogenous spatial occupancy fundamentally alters stochastic hitting times and optimal option exercise thresholds.

Second, the model is deeply informed by the emerging astropolitical literature on the economics of space and orbital debris. [Adilov et al. \(2015\)](#) provide a foundational economic analysis of Earth orbit pollution, establishing that open-access orbital environments naturally

induce a tragedy of the commons, leading to over-launching and excessive debris generation. [Rao et al. \(2020\)](#) quantify this physical tragedy, demonstrating that the implementation of optimal orbital-use fees could more than quadruple the long-term value of the space industry by adequately pricing collision risks. While these papers perfectly diagnose the economic failure of orbital congestion, their proposed policy solutions rely heavily on the existence of a coordinated international governing body capable of levying taxes and enforcing sovereign compliance. I extend this literature by directly addressing the geopolitical reality that international space treaties fail due to sovereign defection. I provide a decentralized alternative, proving that private capital markets and structural financial contracts can substitute for state enforcement to achieve the optimal orbital-use fee.

Third, this research intersects with the broad economic literature on climate change and the governance of the commons. [Ostrom \(1990\)](#) demonstrated how localized institutions can govern common-pool resources without top-down sovereign control. At the macroeconomic level, [Weitzman \(2009\)](#) and [Nordhaus \(2014\)](#) focus on pricing catastrophic climate tipping points and the social cost of carbon. However, traditional climate economics largely views the Earth’s atmosphere as a closed system. By analyzing the transition of heavy artificial intelligence computation to the vacuum of space, this paper introduces a spatial, extraterrestrial dimension to climate economics, where mitigating the thermal constraint of the Earth is achieved through physical atmospheric decoupling rather than mere carbon reduction.

Finally, this paper contributes to the rapidly growing field of sustainable finance and environmental, social, and governance (ESG) contracting. Recent theoretical work by [Pástor et al. \(2021\)](#) has formalized how investor preferences generate a greenium, systematically lowering the cost of capital for sustainable firms. Empirically, [Flammer \(2021\)](#) confirms that the issuance of corporate green bonds yields positive abnormal returns and improves environmental performance, while [Bolton and Kacperczyk \(2021\)](#) show that institutional investors actively demand risk premiums for carbon-intensive assets. However, a significant gap remains in structural corporate finance regarding how these instruments operate dynamically. Much of the literature focuses on the empirical asset pricing of the greenium, lacking a continuous-time framework that shows how sustainability covenants actively alter strategic investment timing among oligopolistic competitors. I contribute to this field by providing a subgame-perfect microfoundation, mathematically proving that firms will voluntarily accept restrictive, long-term environmental covenants not out of altruism, but as a strategic financial weapon to manipulate competitor entry thresholds and extend their own monopoly deterrence windows.

3 Institutional Background

To understand why institutional capital is currently being deployed toward extraterrestrial real estate despite the long-term nature of microwave power transmission, it is critical to deconstruct the immediate, compounding frictions in terrestrial energy markets.

3.1 Terrestrial Frictions: Grid Congestion and the Thermal Constraint

The scaling laws of artificial intelligence dictate that parameter counts and computational requirements grow exponentially. Modern AI training clusters require gigawatts of continuous, uninterrupted baseload power. In the United States, the localized concentration of these hyperscale data centers has broken local grid equilibrium. Empirical observations in regions dense with AI infrastructure, such as Northern Virginia and specific counties in the Pacific Northwest, demonstrate that the massive, inflexible draw of computation centers has triggered systemic grid congestion. As utility providers are forced to initiate rapid, unplanned grid upgrades and source peaker-plant energy, the associated costs are socialized. Recent municipal data indicates that in certain heavily concentrated AI data center zones, local residential and commercial electricity fees have surged by up to 267%.

Beyond simple electricity pricing, AI computations face a severe physical limitation: the thermal constraint. Up to 40% of the energy consumed by a terrestrial data center is utilized purely for cooling systems (e.g., HVAC and liquid cooling loops). Furthermore, the heat dissipated into the Earth’s atmosphere actively accelerates global warming. Even if a terrestrial data center is powered entirely by localized zero-carbon solar panels or nuclear small modular reactors (SMRs), the aggregate thermal exhaust generated by exaflops of computation represents a net-positive heat addition to the fragile biosphere.

3.2 Comparing Alternative Energy Architectures

Technology companies are currently evaluating three primary architectures to solve the AI baseload power problem. First, terrestrial renewables (solar and wind) are environmentally clean but inherently intermittent. They suffer from the duck curve of generation and require massive, prohibitively expensive battery infrastructure to provide the 24/7 reliability AI requires. They are also incredibly land-intensive, triggering heavy local NIMBY (Not In My Backyard) regulatory resistance.

Second, nuclear power (including small modular reactors) offers a highly reliable baseload that can theoretically be deployed much faster than extraterrestrial energy infrastructure.

However, its primary drawback is fundamental to the biosphere: while nuclear power is carbon-free, the massive thermal exhaust generated by gigawatt-scale AI computation remains trapped on Earth. Even if powered by nuclear energy, the sheer volume of heat dissipated by AI data centers directly exacerbates localized and global thermal warming limits.

The third alternative is space-based energy generation. Space solar power currently faces immense drawbacks—specifically astronomical launch costs and highly immature microwave-beaming technology, cementing it as a long-term goal. However, it offers an unparalleled, structural advantage that terrestrial nuclear cannot match: by placing the solar array and the AI data center in orbit, it leverages the vacuum of space as a costless, infinite heat sink. It removes the thermal exhaust of computation entirely from the Earth’s fragile atmosphere, solving the ultimate physical bottleneck of artificial intelligence.

3.3 The Pivot to Space AI Data Centers and the Corporate Land Grab

The historical barrier to space-based solar power has been the physics of the downlink: converting DC power to microwaves, beaming it through the atmosphere, and recapturing it via terrestrial rectennas involves massive transmission losses. Rather than transmitting power to Earth, the optimal architecture is to transmit the data centers to space.

By co-locating the AI computation hardware with the space-based solar array, the energy is consumed locally in orbit with near-100% efficiency. Furthermore, the ambient temperature of deep space serves as an infinite, costless thermal heat sink, instantly eliminating the 40% power tax required for terrestrial cooling and permanently removing computational heat from the Earth’s biosphere. The finished product—trained AI weights or processed inference data—can be transmitted back to Earth via laser optical communications, a technology that is already highly mature and widely deployed.

Because specific orbits (such as Sun-synchronous LEO) offer the optimal balance of continuous sunlight and low latency, they represent finite real estate. Consequently, a massive corporate land grab is actively underway. For instance, SpaceX has submitted regulatory filings detailing plans for an unprecedented 1 million satellite constellation, explicitly targeting advanced orbital data routing and computation. Similarly, Amazon’s Blue Origin has filed Federal Communications Commission (FCC) applications for Project Sunrise, a proposed constellation of up to 51,600 satellites in LEO scheduled to begin deployment in Q4 2027, explicitly designed to establish space-based data center networks.

3.4 The Astropolitical Reality: Why Centralized Treaties Fail

As firms rush to secure these slots, they collide with an anarchic geopolitical reality. The foundational legal framework, the 1967 Outer Space Treaty, strictly prohibits sovereign appropriation of space. There is no central global sovereign capable of auctioning orbital slots, levying Pigouvian taxes, or enforcing environmental quotas. If the United States attempts to unilaterally restrict its domestic corporations to prevent orbital congestion, adversarial nations will simply accelerate their own state-sponsored deployments to capture AI supremacy, rendering domestic regulation worse than useless.

Because sovereign enforcement is non-credible, the preservation of the orbital commons must be achieved through self-enforcing, decentralized financial contracts. Capital markets—specifically institutional debt providers, ESG funds, and insurance syndicates—act as the globally binding enforcement mechanism. If the financing structure is designed correctly, private capital can mathematically engineer the optimal entry delays that international diplomacy cannot.

4 Model Environment and the Standard Equilibrium

I model the investment decision as a dynamic game within a three-regime market lifecycle. The state of the market depends on the number of active firms, denoted as $N_t \in \{0, 1, 2\}$, and a potential terminal congestion shock. The stochastic terrestrial energy price P_t , representing the demand for computation, follows a geometric Brownian motion:

$$dP_t = \alpha P_t dt + \sigma P_t dW_t \tag{1}$$

where α is the drift parameter representing secular growth in demand, σ is the volatility parameter representing uncertainty, and dW_t is the increment of a standard Wiener process. The constant risk-free interest rate is r , and I assume $r > \alpha$ to ensure the present value of future cash flows converges to a finite number. Firms face a sunk capital cost I to deploy extraterrestrial energy infrastructure.

4.1 The Three-Regime Market Lifecycle and Endogenous Risk

Unlike traditional capacity expansion models where exogenous shocks arrive uniformly, the risk of a terminal congestion collapse in this model is strictly endogenous to total spatial occupancy. The market transitions sequentially through three distinct physical regimes.

Regime 1 is the scarcity phase where $N_t = 1$. The first firm, defined as the Leader,

enters the market. Operating without spatial competition, it captures a monopoly cash flow multiplier of $A_{mono}^{std}P_t$. Because a single firm operates with largely unrestricted orbital slots, it can optimize its own navigation and positioning algorithms without external interference. Therefore, it generates only a minimal baseline risk of a self-induced orbital cascade. I model this risk as a Poisson process with a very low arrival rate λ_1 .

Regime 2 is the abundance phase where $N_t = 2$. The second firm, defined as the Follower, enters the market. Competition for prime orbital real estate, alongside standard energy market share competition, drives down profit margins. This yields a lower cash flow multiplier $A_{duo}^{std}P_t$ for both active firms, where $A_{duo}^{std} < A_{mono}^{std}$. Crucially, the total spatial occupancy doubles. The physical crowding and the inability to perfectly coordinate independent, proprietary navigation algorithms cause the risk of a terminal Kessler Syndrome collapse to spike exponentially. The Poisson arrival rate jumps to λ_2 , where $\lambda_2 \gg \lambda_1$. This jump in risk represents the physical manifestation of the tragedy of the commons, where individual rational entry leads to collective physical ruin.

Regime 3 is the congestion collapse. The orbital environment becomes physically unstable due to a Kessler Syndrome event. Upon the arrival of the shock, which occurs at rate λ_1 during the monopoly phase or λ_2 during the duopoly phase, operational efficiency is permanently impaired. The revenue parameter collapses to A_{coll}^{std} , which is strictly less than A_{duo}^{std} .

To establish the baseline of the model, I first derive the value of an active firm strictly in the terminal collapse regime. Because there are no further state transitions after a collapse, the value of the firm is simply the expected present value of a perpetual cash flow growing at rate α . Applying the standard Gordon Growth continuous-time integral for a geometric Brownian motion yields the collapsed project value:

$$V_{coll}^{std}(P_t) = \int_0^\infty \mathbb{E}[A_{coll}^{std}P_{t+\tau}]e^{-r\tau}d\tau = \frac{A_{coll}^{std}P_t}{r - \alpha} \quad (2)$$

4.2 Derivation of Risk-Adjusted Regime Values

Consider the market where firms utilize standard finance without any sustainability covenants. To find the value of a firm operating prior to collapse, I must establish intertemporal arbitrage conditions for both the scarcity and abundance regimes.

By the no-arbitrage principle, the required return on the asset over a differential time step dt must exactly equal the cash flow generated plus the expected capital gain. Let

$\lambda_i \in \{\lambda_1, \lambda_2\}$ be the state-dependent risk. The Hamilton-Jacobi-Bellman (HJB) equation is:

$$rV_{regime}^{std}(P_t)dt = A_{regime}^{std}P_tdt + \mathbb{E}[dV_{regime}^{std}(P_t)] \quad (3)$$

Because the asset is subject to both continuous price diffusion and discrete collapse jumps, I apply Ito's Lemma for jump-diffusion processes to expand the expected change term. The continuous geometric Brownian motion component is expanded using a second-order Taylor series. Simultaneously, the discrete jump component represents the probability $\lambda_i dt$ that the firm suffers a discrete capital loss, falling from the active regime value down to the collapsed state. The total expected differential is:

$$\mathbb{E}[dV_{regime}^{std}(P_t)] = \left(\alpha P_t \frac{\partial V}{\partial P} + \frac{1}{2} \sigma^2 P_t^2 \frac{\partial^2 V}{\partial P^2} \right) dt + \lambda_i (V_{coll}^{std}(P_t) - V_{regime}^{std}(P_t)) dt \quad (4)$$

Substituting this expanded expectation back into the HJB equation and dividing the entire equation by dt yields the fundamental non-homogeneous linear differential equation:

$$\frac{1}{2} \sigma^2 P_t^2 (V_{regime}^{std})'' + \alpha P_t (V_{regime}^{std})' - (r + \lambda_i) V_{regime}^{std} + A_{regime}^{std} P_t + \lambda_i V_{coll}^{std}(P_t) = 0 \quad (5)$$

To solve this differential equation, I guess a particular solution of the form $V = kP_t$ because the underlying cash flows are strictly linear in P_t . Taking the first derivative $V' = k$ and the second derivative $V'' = 0$, I substitute these directly into the differential equation:

$$\frac{1}{2} \sigma^2 P_t^2 (0) + \alpha P_t (k) - (r + \lambda_i) (kP_t) + A_{regime}^{std} P_t + \lambda_i \left(\frac{A_{coll}^{std} P_t}{r - \alpha} \right) = 0 \quad (6)$$

Factoring out P_t from every term (since $P_t > 0$) leaves a purely algebraic equation for k :

$$\alpha k - rk - \lambda_i k + A_{regime}^{std} + \frac{\lambda_i A_{coll}^{std}}{r - \alpha} = 0 \quad (7)$$

Isolating k on the left side yields:

$$k(r + \lambda_i - \alpha) = A_{regime}^{std} + \frac{\lambda_i A_{coll}^{std}}{r - \alpha} \quad (8)$$

Dividing by $(r + \lambda_i - \alpha)$ and multiplying by P_t yields the fundamental, closed-form present value of the active firm in any regime subject to spatial risk λ_i :

$$V_{risky}^{std}(P_t, \lambda_i, A_{regime}^{std}) = \frac{A_{regime}^{std} P_t}{r + \lambda_i - \alpha} + \frac{\lambda_i}{r + \lambda_i - \alpha} \left(\frac{A_{coll}^{std} P_t}{r - \alpha} \right) \quad (9)$$

This derivation explicitly proves the economic intuition of the model. Higher spatial occupancy, represented by the shift from λ_1 to λ_2 , directly cannibalizes the present value of the firm by exponentially increasing the discount factor in the denominator of the primary cash flows.

4.3 The Follower's Option and Threshold Derivation

Before entering the market, the Follower holds a perpetual American call option to invest. Because the option itself yields no cash flows prior to exercise, its value $F(P_t)$ satisfies the standard homogeneous differential equation $0.5\sigma^2 P_t^2 F'' + \alpha P_t F' - rF = 0$. Ruling out asset bubbles and noting that the option value must approach zero as the energy price approaches zero, the standard solution is $F(P_t) = cP_t^\beta$, where $\beta > 1$ is the positive root of the characteristic quadratic equation $0.5\sigma^2\beta(\beta - 1) + \alpha\beta - r = 0$.

The Follower optimizes its entry timing by selecting an optimal threshold P_F^* . Upon entry, the market shifts to the abundance phase, triggering the higher spatial risk λ_2 . At the optimal threshold, two economic boundary conditions must hold. First, value-matching requires that the option value exactly equals the net payoff of investment. Second, smooth-pasting ensures marginal optimization, meaning the derivative of the option equals the derivative of the payoff. The system of equations is:

$$c(P_F^*)^\beta = V_{risk}^{std}(P_F^*, \lambda_2, A_{duo}^{std}) - I \quad (10)$$

$$c\beta(P_F^*)^{\beta-1} = (V_{risk}^{std})'(P_F^*, \lambda_2, A_{duo}^{std}) \quad (11)$$

To solve this algebraically, I isolate the option multiplier c from the smooth-pasting condition:

$$c = \frac{(V_{risk}^{std})'(P_F^*)}{\beta(P_F^*)^{\beta-1}} \quad (12)$$

I then substitute this expression for c directly back into the value-matching condition:

$$\left[\frac{(V_{risk}^{std})'(P_F^*)}{\beta(P_F^*)^{\beta-1}} \right] (P_F^*)^\beta = V_{risk}^{std}(P_F^*) - I \quad (13)$$

Because $\frac{(P_F^*)^\beta}{(P_F^*)^{\beta-1}} = P_F^*$, the equation simplifies fundamentally to:

$$\frac{P_F^* \cdot (V_{risk}^{std})'(P_F^*)}{\beta} = V_{risk}^{std}(P_F^*) - I \quad (14)$$

Recall from Equation 7 that the active firm value is strictly linear in P_t , meaning $V_{risk}^{std}(P_F^*) =$

$P_F^* \cdot (V_{risky}^{std})'(1)$. Therefore, $P_F^* \cdot (V_{risky}^{std})'(P_F^*)$ perfectly equates to $V_{risky}^{std}(P_F^*)$. Substituting this property into the equation yields:

$$\frac{V_{risky}^{std}(P_F^*)}{\beta} = V_{risky}^{std}(P_F^*) - I \quad (15)$$

Rearranging to isolate the project value gives $V_{risky}^{std}(P_F^*) \left(1 - \frac{1}{\beta}\right) = I$, which simplifies to $V_{risky}^{std}(P_F^*) = \frac{\beta}{\beta-1}I$. Finally, dividing by the linear slope $(V_{risky}^{std})'(1)$ yields the Follower's exact standard entry threshold:

$$P_F^* = \frac{\beta}{\beta-1} \frac{I}{(V_{risky}^{std})'(1, \lambda_2, A_{duo}^{std})} \quad (16)$$

4.4 The Leader's Pre-Emptive Race and Value Destruction

A firm contemplating entering the market as the Leader must evaluate the continuous threat of pre-emption. Its post-investment project value, $V_L(P_t)$, equals the value of a perpetual monopoly operating under the low risk λ_1 , minus the expected present value of the competitive loss triggered at the first hitting time when the price reaches the Follower's threshold.

For a geometric Brownian motion starting at P_t , the expected stochastic discount factor evaluated at the random hitting time τ when it reaches an upper boundary P_F^* is mathematically defined as $\mathbb{E}[e^{-r\tau}] = \left(\frac{P_t}{P_F^*}\right)^\beta$. Applying this stochastic discount factor to the discrete loss in value caused by the Follower's entry yields the Leader's value:

$$V_L(P_t) = V_{risky}^{std}(P_t, \lambda_1, A_{mono}^{std}) - \left(\frac{P_t}{P_F^*}\right)^\beta [V_{risky}^{std}(P_F^*, \lambda_1, A_{mono}^{std}) - V_{risky}^{std}(P_F^*, \lambda_2, A_{duo}^{std})] \quad (17)$$

The net present value of leading is defined as $L(P_t) = V_L(P_t) - I$. In a theoretical market where a firm holds an exogenous, unthreatened monopoly right, the firm would maximize its value by finding the absolute peak of the firm value curve. This point is called the socially optimal unthreatened threshold, representing the exact point at which a single firm, facing no threat of competitive entry, would optimally choose to invest. At this threshold, the marginal value of waiting for higher energy prices is exactly zero.

However, in standard market conditions, firms are symmetric and lack exclusive extraction rights. Because each firm holds an identical option to lead and fears being relegated to the less profitable Follower role, they engage in a symmetric pre-emptive race. They will continually deploy capital earlier to undercut their rival. Mathematically, this competitive race terminates strictly at the equilibrium threshold P_L^* where the economic rent of being

the Leader is completely dissipated down to the option value of following. This is defined by the absolute Value-Matching condition:

$$L(P_L^*) = F(P_L^*) \implies V_L(P_L^*) - I = [V_{risky}^{std}(P_F^*, \lambda_2, A_{duo}^{std}) - I] \left(\frac{P_L^*}{P_F^*} \right)^\beta \quad (18)$$

Because the option to follow $F(P)$ is strictly positive and monotonically increasing, setting $L(P_L^*) = F(P_L^*)$ forces the intersection to occur strictly on the upward-sloping portion of the Leader's value curve. This mathematical property guarantees that the pre-emptive threshold is significantly lower than the monopoly optimum ($P_L^* \ll P_M^*$). The standard market yields a Pareto-inefficient race, deploying capital prematurely when market demand is too low to justify the expenditure, severely destroying expected total surplus in a tragedy of the commons.

5 Mechanism Design: Greenwashing vs. Pigouvian Delay

To correct this market failure without a central sovereign, I design the orbital-sustainability bond. This mechanism utilizes two distinct levers to alter the strategic landscape: a clean energy premium and an asymmetric greenium.

A greenium, or green premium, is the reduction in the cost of debt or cost of capital that institutional investors offer to firms that commit to strict, verifiable environmental covenants. Because sustainable funds possess mandates to allocate capital toward climate solutions, they accept slightly lower financial yields in exchange for quantifiable environmental impact, effectively lowering the upfront capital expenditure for the issuing firm.

If a firm issues the bond, verified spatial sustainability allows the firm to sign premium power purchase agreements with consumers, yielding elevated revenue parameters $A^{osb} > A^{std}$. A crucial condition is embedded in the contract: if an orbital congestion collapse occurs, the sustainability certification is immediately revoked. Consequently, the collapsed revenue parameter reverts to the standard baseline ($A_{coll}^{osb} = A_{coll}^{std}$). Consumers will not pay a premium for energy generated from a degraded orbit. In exchange for the upfront financial benefits, both firms bind themselves to a heavy congestion penalty C_p , which is diverted to a global equity fund upon collapse.

The global equity fund is a theoretical contractual trust designed to capture the congestion penalties levied upon firms that trigger orbital degradation. Instead of allowing the surplus to be lost to deadweight inefficiency, this fund contractually redistributes the capital

to terrestrial climate adaptation projects. Unlike purely theoretical welfare measures, the global equity fund represents an actual physical transfer of capital from the extractive firms to global society, thereby fulfilling the Kaldor-Hicks efficiency criterion.

Because the congestion penalty C_p is a fixed nominal liability paid at the random future time of collapse, its expected present value at the moment of entry must be calculated. The time of collapse, τ , is a random variable that follows an exponential distribution with a probability density function $f(t) = \lambda_i e^{-\lambda_i t}$, where λ_i is the state-dependent Poisson arrival rate. Furthermore, I introduce a parameter $\theta \in [0, 1)$ to represent a corporate tax shield or insurance recovery factor on the penalty. In many jurisdictions, catastrophic operational losses or environmental fines can be partially tax-deductible, or a portion may be recoverable through specialized insurance contracts. Therefore, the net penalty paid by the firm is $(1 - \theta)C_p$.

The expected present value of this net penalty for a firm operating in regime i is derived by integrating the continuously discounted net penalty over the exponential probability density of the collapse event:

$$\mathbb{E}[PV(C_{pi})] = \int_0^\infty e^{-rt} [(1 - \theta)C_p] f(t) dt \quad (19)$$

$$= \int_0^\infty e^{-rt} (1 - \theta)C_p \lambda_i e^{-\lambda_i t} dt \quad (20)$$

The constants $(1 - \theta)C_p \lambda_i$ can be factored out of the integral, and the exponential terms can be combined:

$$\mathbb{E}[PV(C_{pi})] = (1 - \theta)C_p \lambda_i \int_0^\infty e^{-(r+\lambda_i)t} dt \quad (21)$$

Evaluating the definite integral yields:

$$\int_0^\infty e^{-(r+\lambda_i)t} dt = \left[-\frac{1}{r + \lambda_i} e^{-(r+\lambda_i)t} \right]_0^\infty = \left(0 - \left(-\frac{1}{r + \lambda_i} \right) \right) = \frac{1}{r + \lambda_i} \quad (22)$$

Substituting this result back yields the closed-form solution for the expected present value of the penalty:

$$\mathbb{E}[PV(C_{pi})] = \frac{\lambda_i}{r + \lambda_i} (1 - \theta)C_p \quad (23)$$

This derivation makes clear that the expected penalty cost is increasing in the physical risk λ_i but decreasing in the financial recovery factor θ .

5.1 The Mathematical Failure of Sovereign Treaties

Before demonstrating the efficacy of private financial mechanisms, it is necessary to mathematically prove why traditional macroeconomic solutions—such as global Pigouvian taxes or international space treaties—structurally fail in astropolitical environments.

A Pigouvian tax is a tax levied on any market activity that generates negative externalities, where the tax is set equal to the social cost of those externalities. In our context, the externality is the marginal increase in collision risk that a second entrant imposes on the first firm and the entire orbital commons. An ideal global sovereign would impose a Pigouvian tax on every satellite launch equal to the expected present value of the damage caused by that satellite’s contribution to orbital congestion. However, the collection of such a tax requires a global entity with unquestioned enforcement power, which does not exist in space.

Lacking a global tax authority, nations might attempt to form a treaty to self-regulate. Consider a treaty where two competing nation-states, State A and State B, agree to delay orbital deployment until a high strategic value threshold is met, effectively agreeing to pay an opportunity cost tax τ by waiting. Entering the orbit early yields a strategic AI supremacy payoff of S . However, because there is no extraterrestrial police force, compliance is entirely voluntary and non-credible.

I model the decision to adhere to the treaty as an optimal stopping problem with defection intensity. Let γ be the Poisson arrival rate of a geopolitical shock that incentivizes immediate defection. This parameter, γ , represents the perceived risk that the rival nation will unilaterally disobey the treaty to seize a first-mover advantage. A higher γ signifies a less trustworthy rival or a more intense geopolitical “arms race” for AI supremacy. If State A complies and waits, its expected value is continuously eroded by the risk that State B will defect, seizing the entire monopoly payoff S for itself and leaving State A with nothing.

The Hamilton-Jacobi-Bellman (HJB) equation for a compliant State A, which holds the option to defect, can be expressed intuitively. The required return on holding the “compliance option” must equal its expected capital appreciation plus its expected loss from being preempted.

$$rV_A(S_t)dt = \mathbb{E}[dV_A(S_t)] + \gamma(0 - V_A(S_t))dt \quad (24)$$

Here, $rV_A(S_t)dt$ is the required return for waiting a moment dt . $\mathbb{E}[dV_A(S_t)]$ is the expected capital gain on the option’s value from the normal stochastic drift of S_t . The crucial term is $\gamma(0 - V_A(S_t))dt$. This is the expected capital loss from geopolitical preemption. With probability γdt , the rival state defects. This event instantly “kills” State A’s option, driving its value to zero. The capital loss is therefore $(0 - V_A(S_t))$.

Expanding the expectation using Ito's Lemma yields the ordinary differential equation:

$$\frac{1}{2}\sigma^2 S^2 V_A'' + \alpha S V_A' - (r + \gamma)V_A = 0 \quad (25)$$

Solving this homogeneous differential equation shows that the effective discount rate on compliance is $r + \gamma$. A higher defection risk γ dramatically increases the cost of waiting, forcing State A to lower its own defection threshold. As $\gamma \rightarrow \infty$ in an intense AI arms race, the compliance threshold collapses toward zero. Since the payoff of defecting ($V_{defect} = S_t$) is always greater than the payoff of complying and paying the tax ($V_{comply} = S_t - \tau$), defection is a dominant strategy. The international treaty is an empty threat, necessitating the creation of the orbital-sustainability bond, which is enforced by capital markets, not states.

5.2 The Mathematical Definition of Greenwashing

The introduction of sustainable finance does not automatically resolve the tragedy of the commons. If the financial contract is poorly calibrated, it can paradoxically accelerate the very problem it is designed to solve. This phenomenon is known as greenwashing. In the context of this model, greenwashing is a state where a firm can issue a "green" or "sustainable" bond, reaping the financial rewards, without being subject to a penalty strong enough to deter socially harmful behavior. In this dual failure, both the Leader and the Follower are capable of issuing such bonds.

Mathematically, greenwashing is defined as a parameterization where the financial benefits of the bond for the secondary entrant—specifically the clean energy revenue premiums (A_{duo}^{osb}) and any potential capital cost reductions—strictly exceed the expected present value of the environmental penalty ($\mathbb{E}[PV(C_{p2})]$). For example, imagine a bond provides an upfront benefit worth \$100 million in net present value, but the expected penalty for causing congestion is only \$10 million. The firm has a net incentive of \$90 million to enter early and congest the orbit.

Under these parameters, the bond acts as a net financial subsidy that inflates the asset's value without deterring risky behavior. This causes the denominator of the Follower's entry threshold equation to expand faster than the numerator. Consequently, the Follower is prompted to enter at a lower price threshold than in the standard unregulated market ($\tilde{P}_F^* < P_F^*$). By financially incentivizing earlier secondary entry, the poorly calibrated green bond accelerates the transition to the high-risk duopoly phase (λ_2), directly accelerating the physical degradation of the shared orbital resource. It is a "green" instrument in name only, while being destructive in practice.

5.3 The Asymmetric Greenium and Pigouvian Delay

To prevent greenwashing and restore efficiency, the optimally designed mechanism relies on a fundamental asymmetry in sustainable finance. ESG institutional investors rationally allocate greenium capital exclusively to the pioneer who establishes the essential clean infrastructure, represented by the Leader, such that $g_L > 0$. They strictly deny the greenium to the Follower ($g_F = 0$) because the secondary entry provides redundant capacity while actively spiking the spatial collapse risk.

The Follower, facing no greenium, evaluates its new entry threshold under the bond covenant. By calibrating congestion penalty C_p to be massive, the mechanism designer inflates the Follower's fixed cost barrier. The numerator in the threshold equation, which includes the sunk cost and the expected penalty, strictly dominates the enhanced revenue multiplier in the denominator. This elevates the Follower's hurdle rate, imposing a Pigouvian delay ($\tilde{P}_F^* > P_F^*$). The mechanism ensures the orbit is only crowded when market demand is extraordinarily high.

A critical theoretical hurdle is explaining why a cost-disadvantaged Follower would voluntarily adopt a high-penalty bond. The Follower chooses the bond framework over standard finance because consumers offer a significant price premium A_{duo}^{osb} for certified energy. At the moment of entry $\tilde{P}_F^*$, the perpetuity value of this premium completely overwhelms the one-time fixed penalty:

$$V_{risky}^{osb}(\tilde{P}_F^*, \lambda_2) - I - \mathbb{E}[PV(C_{p2})] > V_{risky}^{std}(\tilde{P}_F^*, \lambda_2) - I \quad (26)$$

Because this condition holds strictly, the Follower rationally restricts itself to the mechanism.

5.4 Proof of Strategic Dominance and the Double Dividend

The Leader strictly prefers the optimal bond because it acts as a strategic moat. Because the solo-risk λ_1 is negligible, the expected penalty poses almost zero upfront cost to the Leader ($\mathbb{E}[PV(C_{p1})] \approx 0$).

Crucially, the asymmetric greenium mathematically halts the destructive pre-emptive race. The logic of the equilibrium rests on the strategic interaction between the advantaged firm (Firm A, which receives the greenium g_L) and the disadvantaged firm (Firm B, which does not). In a symmetric pre-emptive race, firms invest earlier and earlier until the value of being the leader is driven down to the value of being the follower. This defines the pre-emption boundary.

Let Firm B, the disadvantaged Follower, set this boundary. Firm B will be willing to

pre-empt the market down to a threshold $P_{L,B}^*$ where its net present value of leading, $L_B(P)$, is competed down to its option value of following, $F_B(P)$. At this point, Firm B is indifferent between leading and following. Firm A, the advantaged Leader, knows this. To win the race, Firm A only needs to pre-empt by an infinitesimally small margin, investing at the exact same threshold $P_{L,B}^*$.

At this price, Firm A receives the same project value that Firm B would have received, which is by definition equal to $F_B(P_{L,B}^*)$. However, Firm A has an additional, private benefit that Firm B does not: its effective capital cost is lower by the amount $I \cdot g_L$ due to the greenium. This cost advantage is pure economic rent that is not competed away. Therefore, Firm A's total net present value at the moment of entry is:

$$L_A(P_{L,B}^*) = F_B(P_{L,B}^*) + I \cdot g_L \quad (27)$$

This equation proves that Firm A captures a protected rent equal to the full value of its cost advantage. This generates the double dividend. The greenium g_L shifts the Leader's entry threshold lower, accelerating the initial clean energy transition. Simultaneously, the heavy penalty C_p delays the Follower, significantly extending the Leader's monopoly phase. This combination ensures the Leader's expected value surges, making the bond a strictly dominant strategy.

5.5 Mechanism Implementability: Escrow, IR, and IC Constraints

For this market mechanism to be structurally valid in an anarchic astropolitical environment, it must be mathematically self-enforcing. Because there is no extraterrestrial sovereign to collect the penalty C_p after a Kessler Syndrome event, a post-collapse collection effort would be blocked by corporate bankruptcy shields. Therefore, the mechanism is only implementable if it utilizes an ex-ante financial escrow.

At the time of bond issuance, a fraction of the raised capital must be locked in a trust (or a decentralized cryptographic smart contract) equal to the penalty magnitude $K = C_p$. This trust is governed by an automated oracle linked to European Space Agency (ESA) conjunction data. If a catastrophic collision is verified, the escrow is automatically transferred to the global equity fund.

To prove this escrow mechanism is theoretically sound, it must strictly satisfy the two fundamental axioms of mechanism design: Individual Rationality (IR) and Incentive Compatibility (IC).

1. Individual Rationality (IR): The firm must voluntarily prefer to enter the restrictive escrow contract rather than utilizing standard, unconstrained financing. Let K be the

locked escrow. The expected present value of the escrow loss is $\mathbb{E}[PV(K)] = \frac{\lambda_i}{r+\lambda_i}K$. The IR constraint requires that the net value of the firm under the orbital-sustainability bond (with the greenium g_L and baseload premium A_{mono}^{osb}) exceeds the standard unregulated value:

$$V_L^{osb}(P_L^*) - I - \frac{\lambda_1}{r_L + \lambda_1}K \geq V_L^{std}(P_L^*) - I \quad (28)$$

Because $A_{mono}^{osb} > A_{mono}^{std}$ and $r_L < r$, the left side of the inequality expands significantly faster than the right side. By calibrating K such that the inequality binds with a strict positive surplus for the leader, the IR constraint is satisfied. The pioneer firm willingly posts the collateral.

2. Incentive Compatibility (IC): Once the escrow is posted, the firm must not have a financial incentive to secretly deviate to a high-risk, cost-cutting operational trajectory (e.g., refusing to spend fuel on collision avoidance). Let x be the unobservable safety effort exerted by the firm, costing $c(x)$, which reduces the collision risk $\lambda(x)$. The firm's payoff under the escrow contract is $V_L^{osb} - c(x) - \frac{\lambda(x)}{r_L + \lambda(x)}K$. The IC constraint mandates that the firm's private optimization problem naturally selects the socially optimal high safety effort x^* over a zero-effort defection $x = 0$:

$$V_L^{osb} - c(x^*) - \frac{\lambda(x^*)}{r_L + \lambda(x^*)}K \geq V_L^{osb} - 0 - \frac{\lambda(0)}{r_L + \lambda(0)}K \quad (29)$$

Rearranging this inequality isolates the role of the escrow:

$$K \left(\frac{\lambda(0)}{r_L + \lambda(0)} - \frac{\lambda(x^*)}{r_L + \lambda(x^*)} \right) \geq c(x^*) \quad (30)$$

This derivation proves the exact required size of the penalty. The escrow K must be set precisely high enough that the expected marginal savings in retained collateral strictly exceed the marginal cost of the safety effort $c(x^*)$. By satisfying both the IR and IC constraints, this framework proves that private capital markets can mathematically perfectly replicate the enforcement power of a sovereign state.

6 Spatial Arbitrage and Extraterrestrial AI Data Centers

The baseline model assumes energy is beamed back to the terrestrial grid. However, given the technological friction of microwave transmission, this section analyzes a strategic pivot: deploying AI data centers directly to space. This transforms the model into a problem of

spatial arbitrage.

Deploying AI compute to deep space (e.g., Lagrange points) resolves the congestion risk ($\lambda_{L1} = 0$). However, it introduces two frictions: higher capital expenditure ($I_{L1} > I_{GEO}$) and severe power transmission loss. To model this loss, I introduce $\tau \in [0, 1]$ as the power transmission efficiency coefficient. It represents the fraction of power generated in space that successfully reaches the terrestrial grid. For deep space, this is very low ($\tau \rightarrow 0$).

This technological reality motivates a change in the firm's business model. Instead of selling MWh of power to Earth, the firm's primary product becomes selling TFLOPS of AI computation directly in space. The firm now has two potential revenue streams: (1) a primary, high-growth stream from selling in-space AI compute services, whose demand is modeled by a stochastic process D_t , and (2) a secondary, highly inefficient stream from beaming any marginal, unused power back to Earth, generating revenue $\tau \cdot A_{mono}^{osb} P_t$. This dual-revenue structure explicitly captures the firm's strategic choice to prioritize the high-margin data market over the low-efficiency energy market.

To circumvent this transmission loss, the firm exercises a technological pivot. The solar array powers the supercomputers locally, and the data center utilizes the ambient vacuum of space for costless thermal cooling. The product transmitted back to Earth is processed computational data via lasers, suffering negligible loss. I model the demand for this space-based computation as an independent geometric Brownian motion D_t , evolving with drift α_D and volatility σ_D :

$$dD_t = \alpha_D D_t dt + \sigma_D D_t dW_t^D \quad (31)$$

To rigorously derive the value of a firm deploying an integrated AI center to deep space, $V_{L1}(P_t, D_t)$, we construct the bivariate expected present value integral. The expectation operator under the risk-neutral measure is:

$$V_{L1}(P_0, D_0) = \mathbb{E}_0 \left[\int_0^\infty e^{-rt} (\tau \cdot A_{mono}^{osb} P_t + D_t) dt \right] \quad (32)$$

Because P_t and D_t are governed by independent geometric Brownian motions, the expectation of a sum is the sum of expectations. We separate the integral:

$$V_{L1}(P_0, D_0) = \int_0^\infty e^{-rt} \tau A_{mono}^{osb} \mathbb{E}_0[P_t] dt + \int_0^\infty e^{-rt} \mathbb{E}_0[D_t] dt \quad (33)$$

Using the standard property of geometric Brownian motion, the expected value at time t

given initial value X_0 is $\mathbb{E}_0[X_t] = X_0 e^{\alpha t}$. Substituting this into both integrals:

$$V_{L1}(P_0, D_0) = \int_0^\infty \tau A_{mono}^{osb} P_0 e^{-(r-\alpha)t} dt + \int_0^\infty D_0 e^{-(r-\alpha_D)t} dt \quad (34)$$

Evaluating the definite integrals from 0 to ∞ (assuming $r > \alpha$ and $r > \alpha_D$ to satisfy the transversality condition preventing infinite asset bubbles) yields the closed-form perpetuity value:

$$V_{L1}(P_t, D_t) = \frac{\tau \cdot A_{mono}^{osb} P_t}{r - \alpha} + \frac{D_t}{r - \alpha_D} \quad (35)$$

Intuitively, this equation proves that while the firm loses efficiency on terrestrial power beaming ($\tau \approx 0$ drives the first term down), it gains a perpetual, risk-free exposure to the hyper-growth AI computation economy ($\frac{D_t}{r-\alpha_D}$). Because deep space faces zero congestion risk, there is no jump-to-default discount in the denominator.

The mechanism designer can calibrate the orbital-sustainability bond not merely to delay the follower, but to induce a permanent spatial pivot. A firm facing the disadvantaged follower role evaluates two mutually exclusive net present values: entering near-Earth orbit subject to the bond penalty C_p , or entering deep space subject to the higher rocket payload cost I_{L1} . The spatial pivot occurs at the exact terrestrial threshold P_{pivot} where the two values perfectly equalize:

$$V_{GEO}^{osb}(P_{pivot}, \lambda_2) - I_{GEO} - \mathbb{E}[PV(C_p)] = V_{L1}(P_{pivot}, D_0) - I_{L1} \quad (36)$$

By elevating the near-Earth congestion penalty C_p to a critical threshold, the expected financial cost of terrestrial compliance strictly exceeds the transmission and payload frictions of deep space. The optimal strategy of the follower shifts from crowding near-Earth orbit to establishing autonomous data centers in deep space. This mathematically achieves a Pareto-optimal allocation: it secures terrestrial energy safety, actively utilizes the vacuum of space to halt the thermal warming of the Earth, and forces the secondary firm to lay the computational foundation for the future extraterrestrial economy.

7 Macroeconomic Spillovers: Terrestrial Grid Relief and Social Equity

The analysis thus far has focused on the localized, private corporate finance benefits of extraterrestrial expansion. However, to fully capture the Kaldor-Hicks welfare improvements, the model must be extended to incorporate the macroeconomic spillovers onto the terrestrial

utility grid.

A primary political and economic friction of terrestrial artificial intelligence expansion is the crowding-out effect. Hyperscale data centers possess highly inelastic power demands, meaning they will purchase electricity regardless of the wholesale price. When this massive, inelastic demand is layered onto a rigid terrestrial power grid, it forces utility providers to activate expensive, carbon-intensive peaker plants, driving up the marginal cost of electricity for all local residential and commercial consumers. By designing a financial mechanism that successfully relocates these AI computations to deep space, the model triggers a profound, secondary macroeconomic benefit: Terrestrial Grid Relief.

To mathematically formalize this social equity dividend, consider a regional terrestrial electricity market. The inverse supply curve of the terrestrial grid is strictly convex, reflecting the increasing marginal cost of power generation, modeled as $P(Q) = \xi Q^\eta$, where $\xi > 0$ is a scale parameter and $\eta > 1$ represents the elasticity of generation.

Total initial terrestrial demand is the sum of baseline retail demand Q_{retail} and inelastic AI data center demand Q_{AI} . The pre-transition equilibrium wholesale price is therefore heavily inflated:

$$P_{\text{pre}} = \xi(Q_{\text{retail}} + Q_{\text{AI}})^\eta \quad (37)$$

When the orbital-sustainability bond successfully funds the deep-space AI transition, the tech company exercises its spatial pivot, permanently withdrawing Q_{AI} from the Earth's biosphere. The terrestrial market clears at a new, structurally lower equilibrium price:

$$P_{\text{post}} = \xi(Q_{\text{retail}})^\eta \quad (38)$$

Because $\eta > 1$, the removal of the marginal AI load collapses the price exponentially, yielding $P_{\text{post}} \ll P_{\text{pre}}$.

The exact monetary value of this relief to the terrestrial public is quantified by the change in Consumer Surplus (ΔCS). Integrating the inverse supply function bounds the total savings achieved by regular consumers who no longer have to compete with technology giants for baseline electricity:

$$\Delta CS = \int_0^{Q_{\text{retail}}} (P_{\text{pre}} - P_{\text{post}}) dq \quad (39)$$

$$= Q_{\text{retail}} \cdot [\xi(Q_{\text{retail}} + Q_{\text{AI}})^\eta - \xi(Q_{\text{retail}})^\eta] \quad (40)$$

This mathematical derivation establishes a powerful Triple Dividend inherent to the optimally designed orbital-sustainability bond. First, it provides uninterrupted, zero-carbon

energy to the frontier technology sector. Second, it utilizes the vacuum of space to permanently remove the catastrophic thermal exhaust of artificial intelligence from the Earth's warming atmosphere. Third, and most crucially for public policy, it structurally deflates terrestrial utility bills, generating massive consumer surplus and ensuring affordable, fair energy pricing for human populations on Earth.

This proves that financing the space-based energy transition is not merely an exercise in astropolitical expansion; it is simultaneously one of the most effective terrestrial ESG and social equity investments available to modern capital markets. By utilizing structured private finance to move industrial loads off-planet, capital markets can resolve domestic energy poverty far more effectively than through localized terrestrial subsidies.

8 Endogenous Risk Mitigation: Reusability and Debris Capture

The standard jump-diffusion literature treats failure rates as exogenous constants. In reality, the advent of reusable rockets, autonomous debris-catching mechanisms, and in-orbit servicing allows firms to actively invest capital to suppress environmental risk.

Let $x \geq 0$ denote the discretionary capital allocated to reusability and debris mitigation. The total sunk cost becomes $I(x) = I_0 + x$. The endogenous congestion risk is a convex, decreasing function of this investment, modeled as $\lambda(x) = \bar{\lambda}e^{-\kappa x}$, where $\kappa > 0$ represents the technological efficacy of robotic systems.

At the moment of entry, the firm must solve a static optimization problem to determine the optimal level of risk mitigation x^* . The firm minimizes its total expected friction, defined as the sum of the upfront mitigation cost x and the expected present value of the congestion penalty. Because the collapse arrival follows a Poisson process with rate $\lambda(x)$, the expected present value of the penalty paid at the random time of collapse τ is $\int_0^\infty e^{-r_L t} \lambda(x) e^{-\lambda(x)t} C_p dt$. Evaluating this integral yields the minimization function:

$$\min_{x \geq 0} \left[x + \frac{\lambda(x)}{r_L + \lambda(x)} C_p \right] = \min_{x \geq 0} \left[x + \frac{\bar{\lambda}e^{-\kappa x}}{r_L + \bar{\lambda}e^{-\kappa x}} C_p \right] \quad (41)$$

To find the optimal investment level x^* , I take the first derivative of the friction function with respect to x and set it to zero. Applying the quotient rule $\left(\frac{u}{v}\right)' = \frac{u'v - uv'}{v^2}$ to the penalty fraction, where $u = \bar{\lambda}e^{-\kappa x}$ and $v = r_L + \bar{\lambda}e^{-\kappa x}$, yields $u' = -\kappa\bar{\lambda}e^{-\kappa x}$ and $v' = -\kappa\bar{\lambda}e^{-\kappa x}$. The

derivative of the fraction becomes:

$$\frac{(-\kappa\bar{\lambda}e^{-\kappa x})(r_L + \bar{\lambda}e^{-\kappa x}) - (\bar{\lambda}e^{-\kappa x})(-\kappa\bar{\lambda}e^{-\kappa x})}{(r_L + \bar{\lambda}e^{-\kappa x})^2}C_p = \frac{-r_L\kappa\bar{\lambda}e^{-\kappa x}}{(r_L + \bar{\lambda}e^{-\kappa x})^2}C_p \quad (42)$$

Adding the derivative of the upfront cost (1) yields the final first-order condition:

$$1 - \frac{\kappa\bar{\lambda}e^{-\kappa x}r_L}{(r_L + \bar{\lambda}e^{-\kappa x})^2}C_p = 0 \quad (43)$$

This derivation proves a vital contracting insight: absent the orbital-sustainability bond ($C_p = 0$), the first-order condition collapses to $1 - 0 = 1$. Because the derivative is strictly positive for all x , the cost function is strictly increasing. Therefore, the optimal private investment in reusability and debris capture is cornered precisely at zero ($x^* = 0$). Standard market competition inherently disincentivizes environmental protection because the firm bears the full capital expenditure x but receives zero financial benefit for reducing the global parameter λ .

However, under the bond framework, $C_p > 0$ introduces a heavy financial consequence. The firm endogenously selects an interior optimum $x^* > 0$, perfectly internalizing the externality. The optimal entry threshold P^* is subsequently derived by substituting the minimized cost $I(x^*)$ and optimized risk $\lambda(x^*)$ back into the continuous-time value-matching and smooth-pasting conditions.

9 Financial Architecture: Staged Deployment, SPVs, and Risk Sharing

The baseline mathematical model, while useful for establishing equilibrium dynamics, makes a simplifying assumption common in real options theory: that the entire \$14 billion capital expenditure I is deployed instantaneously as a single, irreversible shock. In reality, financing and deploying a frontier megaproject of this scale is a far more complex corporate finance challenge. No rational board or venture capital syndicate would approve such a massive, monolithic capital outlay given the extreme technological and market uncertainties.

This section, therefore, deconstructs this financing problem and provides a practical, multi-stage financial architecture that makes the project viable for institutional investors and technology firms. The purpose here is to bridge the gap between abstract game theory and actionable corporate strategy. The three subsections are designed to be read sequentially, as they solve three distinct but interconnected layers of the financing problem:

1. Managing Technological and Market Risk (Staged Deployment): The first challenge is managing the immense uncertainty. I address this by replacing the single investment with a compound real option. This allows the firm to first deploy a smaller, less expensive LEO demonstration constellation. This initial stage acts as a pilot program, resolving key technological and market questions before the main capital is committed. I will show how this dramatically lowers the initial barrier to entry.
2. Managing Balance Sheet Risk (SPVs and Tranching): Even if staged, the project is too risky for a technology giant to hold directly on its corporate balance sheet, as it would harm their asset-light valuation multiples. The solution is to create a legally separate Special Purpose Vehicle (SPV). I will detail how this SPV can raise capital by issuing different tranches (slices) of risk—specifically senior debt and subordinated equity—to different classes of investors.
3. Managing Cash Flow and Default Risk (PPA Collars): Finally, the SPV must be able to service its debt. To guarantee its revenue stream and prevent default during periods of low energy prices, the SPV needs a credit enhancement. I model this using a Power Purchase Agreement (PPA) with a derivative collar. I will mathematically prove that a PPA with a price floor is equivalent to an embedded put option that injects synthetic capital into the SPV, securing the debt tranche and making the entire structure bankable.

Together, these three financial structures form a comprehensive and realistic blueprint for financing the extraterrestrial energy transition.

9.1 Compound Real Options and Backward Induction

To formalize this staged financing, I expand the model into a compound real options framework. Let the total required capital be bifurcated such that $I = I_1 + I_2$. Phase 1 (I_1) represents the LEO demonstration, and Phase 2 (I_2) represents the GEO scale-up. Because a LEO demonstration satellite is not continuously illuminated by the sun, it generates a fractional, intermittent cash flow represented by the multiplier A_{demo} , where $A_{demo} < A_{mono}^{osb}$. Upon the successful completion of Phase 1, the firm possesses an active underlying asset yielding A_{demo} alongside a perpetual unexercised real option to upgrade to Phase 2.

This two-stage approach offers profound benefits in risk management and information acquisition, which are critical for securing initial funding. The Phase 1 LEO demonstration serves as a strategic pilot program with several objectives. First, it resolves technological uncertainty by proving the viability of the satellite hardware, solar collection efficiency, and

local power management systems in the harsh environment of space. Second, it resolves market uncertainty by establishing initial contracts and demonstrating the ability to meet the operational demands of a nascent in-space AI compute market.

From a corporate finance perspective, staging the investment fundamentally de-risks the project. If the LEO demonstration fails or reveals that market demand is weaker than anticipated, the firm can abandon the project, limiting its maximum loss to the smaller initial investment I_1 . It is not obligated to deploy the massive follow-on capital I_2 . Conversely, a successful LEO phase provides an invaluable signal to capital markets that the technology is sound and the business model is viable, making it significantly easier and cheaper to raise the larger tranche of capital required for the full GEO scale-up. This sequential resolution of uncertainty is the core economic value of the compound option.

Because this is a sequential investment game, I must solve for the optimal thresholds using backward induction. I first determine the optimal threshold to initiate the Phase 2 upgrade, denoted P_2^* . The firm evaluates the marginal payoff of the upgrade, which is the exact difference in perpetuity values between the GEO baseload capability and the LEO intermittent capability, minus the scale-up cost I_2 . The value of the Phase 2 upgrade option, $F_2(P_t) = c_2 P_t^\beta$, must satisfy the following value-matching and smooth-pasting conditions at the boundary P_2^* :

To solve for the optimal exercise strategy, we must first define the value of the unexercised option to upgrade. The value of this perpetual American-style option, $F_2(P_t)$, must satisfy the standard Hamilton-Jacobi-Bellman (HJB) equation for an asset with no cash flows: $\frac{1}{2}\sigma^2 P^2 F_2'' + \alpha P F_2' - r_L F_2 = 0$. The general solution to this second-order ordinary differential equation is of the form $F_2(P_t) = c_{2,A} P_t^\beta + c_{2,B} P_t^{\beta_N}$, where $\beta > 1$ and $\beta_N < 0$ are the positive and negative roots of the fundamental quadratic equation. However, because the option to upgrade is a call option, its value must approach zero as the underlying energy price P_t approaches zero. The term $P_t^{\beta_N}$ would explode to infinity as $P_t \rightarrow 0$, so we must impose the boundary condition that its coefficient, $c_{2,B}$, is zero. This leaves us with the standard functional form for a perpetual call option: $F_2(P_t) = c_2 P_t^\beta$.

At the yet-to-be-determined optimal exercise threshold P_2^* , two economic conditions must be met to ensure a rational and optimal exercise decision:

1. Value-Matching (Continuity Condition): At the exact moment of exercise, the value of the asset held just before the decision (the option, $F_2(P_2^*)$) must equal the net value of what is received immediately after the decision. If they were not equal, an arbitrage opportunity would exist. The firm receives an upgraded project with a value of $\frac{A_{mono}^{osb} P_2^*}{r_L - \alpha}$ but gives up the existing LEO project with value $\frac{A_{demo} P_2^*}{r_L - \alpha}$ and pays the capital cost I_2 . The net payoff is therefore the difference in project values minus the investment cost.

This gives us the first boundary condition.

2. Smooth-Pasting (Optimality Condition): This is the first-order condition for an optimal exercise boundary. It requires the derivatives of the option value and the payoff function to be equal at the threshold: $F_2'(P_2^*) = (\text{Payoff})'(P_2^*)$. Intuitively, if the two curves did not meet smoothly (i.e., they met at a kink), the holder of the option could improve their outcome by changing the exercise strategy. For example, if the slope of the option value was greater than the slope of the payoff at the boundary, it would imply that the marginal value of waiting an instant longer was still positive, meaning the decision to exercise was premature. A smooth paste ensures that the marginal benefit of waiting is exactly zero at the optimal threshold.

These two conditions provide the following system of equations used to solve for the two unknowns, the constant c_2 and the threshold P_2^* :

$$c_2(P_2^*)^\beta = \frac{A_{mono}^{osb} P_2^*}{r_L - \alpha} - \frac{A_{demo} P_2^*}{r_L - \alpha} - I_2 \quad (44)$$

$$c_2 \beta (P_2^*)^{\beta-1} = \frac{A_{mono}^{osb} - A_{demo}}{r_L - \alpha} \quad (45)$$

To solve this system algebraically, I isolate the option multiplier c_2 from the smooth-pasting condition (Equation 23) and substitute it directly into the value-matching condition (Equation 22). Factoring out P_2^* and rearranging for the boundary yields the Phase 2 execution threshold:

$$P_2^* = \frac{\beta}{\beta - 1} \frac{I_2(r_L - \alpha)}{A_{mono}^{osb} - A_{demo}} \quad (46)$$

Next, rolling backward in time, I evaluate the initial threshold P_1^* required to initiate Phase 1. Before any investment is made, the firm holds a primary option $F_1(P_t) = c_1 P_t^\beta$. Upon investing I_1 at the boundary P_1^* , the firm receives the Phase 1 cash flows plus the newly activated embedded option to upgrade later $F_2(P_t)$. The value-matching condition for Phase 1 is therefore:

$$c_1(P_1^*)^\beta = \frac{A_{demo} P_1^*}{r_L - \alpha} + c_2(P_1^*)^\beta - I_1 \quad (47)$$

This value-matching condition has a clear economic intuition. Upon exercising the first option (valued at $c_1(P_1^*)^\beta$), the firm pays the cost I_1 and receives two distinct assets: (1) the stream of cash flows from the Phase 1 LEO project, valued at $\frac{A_{demo} P_1^*}{r_L - \alpha}$, and (2) the now-active real option to upgrade to Phase 2, valued at $c_2(P_1^*)^\beta$. The equation simply states that the value of the initial option must equal the value of the two assets it creates, net of the initial cost.

By moving the embedded option term $c_2(P_1^*)^\beta$ to the left side of the equation, I can define a net option coefficient $c_{net} = c_1 - c_2$. The net option coefficient has an important economic interpretation. Since the initial option F_1 is an option to acquire both the LEO project and the upgrade option F_2 , the difference, $F_1 - F_2$, can be thought of as the value of an option to acquire only the LEO project's cash flows. This algebraic simplification allows us to treat the problem as if we are solving for a standard real option on a single asset (the LEO project), making the mathematics tractable and leading directly to a clean, closed-form solution for the Phase 1 threshold. The equation simplifies fundamentally to $c_{net}(P_1^*)^\beta = \frac{A_{demo}P_1^*}{r_L - \alpha} - I_1$. Taking the first derivative with respect to P yields the Phase 1 smooth-pasting condition:

$$c_{net}\beta(P_1^*)^{\beta-1} = \frac{A_{demo}}{r_L - \alpha} \quad (48)$$

Dividing the updated value-matching condition by the updated smooth-pasting condition mathematically eliminates the constant c_{net} , yielding the closed-form threshold for Phase 1 entry:

$$P_1^* = \frac{\beta}{\beta - 1} \frac{I_1(r_L - \alpha)}{A_{demo}} \quad (49)$$

This mathematical reduction reveals a highly desirable economic property: the initial barrier to entry P_1^* completely decouples from the massive terminal scale-up cost I_2 . It relies strictly on the fraction of the demonstration cost I_1 and its specific revenue capability A_{demo} . The economic intuition is that the firm has legally decoupled its downside risk; if terrestrial energy demand unexpectedly plummets, the firm can optimally choose to never exercise $F_2(P_t)$, limiting its maximum sunk loss purely to I_1 . By structuring the space infrastructure transition as a compound real option, tech companies can drastically lower the immediate deployment hurdle, accelerating the timeline to launch initial space-based AI capabilities without exposing their corporate balance sheets to the terminal \$14 billion risk.

9.2 SPV Tranching and Jump-to-Default Credit Valuation Adjustment

Even with staged deployment, scaling the primary GEO infrastructure requires external institutional capital. Technology giants rely heavily on asset-light valuations and avoid consolidating unproven heavy infrastructure on their corporate balance sheets. Therefore, the optimal financing architecture is the creation of a strictly ring-fenced Special Purpose Vehicle (SPV) that issues tranching debt and equity.

A tranche refers to a segment of a pooled financial instrument or securitization that has been split up by risk, maturity, or other characteristics. In the context of SPV project finance,

creating tranches allows the issuer to appeal to a wider range of investors with different risk appetites. For this project, I create two primary tranches of the capital structure: a senior debt tranche, which has the first claim on the SPV's cash flows and is therefore lower risk, and a subordinated equity tranche, which has a residual claim on cash flows (i.e., it gets paid only after the debt holders are paid) and is therefore much higher risk. This structure allows risk-averse investors like pension funds and ESG bond funds to participate in the low-risk senior debt, while venture capital and the founding company can take on the high-risk, high-return equity portion.

Let the SPV raise the total capital $I = D + E$, where D represents senior secured debt and E represents subordinated equity. The senior debt is issued as a Green Bond to ESG-focused institutional investors, carrying an initial risk-free rate discounted by the greenium $r_D = r - g$. However, because the SPV's singular asset is located in orbit, it is continuously subject to the endogenous Kessler Syndrome collapse risk λ_1 . Therefore, the debt is not genuinely risk-free. If a terminal orbital collapse occurs, the SPV defaults instantaneously, and the physical asset is destroyed, implying a recovery rate of strictly zero.

To properly price the continuous debt coupon c required by bondholders to yield their target rate r_D , I apply a structural Merton-style credit risk framework incorporating a Poisson jump-to-default intensity. The fundamental value of the debt evaluated under the physical risk measure is the integral of the coupon discounted over time, multiplied by the survival probability $\mathbb{P}(\tau > t)$:

$$D = \int_0^\infty e^{-r_D t} c \cdot \mathbb{P}(\tau > t) dt = \int_0^\infty c \cdot e^{-r_D t} e^{-\lambda_1 t} dt \quad (50)$$

Combining the exponents yields $e^{-(r_D + \lambda_1)t}$. Evaluating the integral from zero to infinity yields the closed-form present value of the debt tranche:

$$D = \frac{c}{r_D + \lambda_1} \quad (51)$$

Rearranging this equation yields the optimal continuous coupon strictly required to float the debt: $c^* = D(r_D + \lambda_1)$. Expanding this yields $c^* = D \cdot r_D + D \cdot \lambda_1$. The secondary term, $\lambda_1 D$, acts as a mathematical Credit Valuation Adjustment (CVA). It represents the exact endogenous risk premium institutional investors mandate to compensate for orbital congestion exposure.

Furthermore, following Modigliani-Miller Proposition II dynamics, the required return on equity r_E is highly sensitive to the chosen leverage ratio $L = D/I$. As leverage increases, the absolute fixed coupon c^* scales linearly. This structurally strips free cash flow away from

the subordinated equity tranche. Because the equity holders now absorb the exact same operational and collapse volatility but receive a significantly thinner slice of the residual cash flows, their financial risk spikes exponentially, forcing r_E fundamentally higher.

9.3 Agency Frictions and the Asset Substitution Problem

While tranching successfully distributes risk and isolates the corporate balance sheet, it introduces a classic corporate finance friction: the agency cost of debt. Following the foundational framework of [Meckling and Jensen \(1976\)](#), the issuance of fixed-claim senior debt inside the SPV fundamentally alters the incentives of the subordinated equity holders.

Because the equity holders possess a residual claim on the SPV, their payoff profile is asymmetric; it is structurally identical to a call option on the underlying astropolitical asset. Once the Green Bond is issued and the capital I is sunk, the equity holders are strictly incentivized to maximize the volatility of the firm and extract maximum near-term cash flows, even if it physically degrades the orbital environment and spikes the terminal collapse risk λ_1 . This is the asset substitution (risk-shifting) problem. The equity holders capture 100% of the upside of aggressive orbital maneuvering, but the downside risk of a Kessler Syndrome collapse—which drives the asset value to zero—is disproportionately absorbed by the institutional bondholders.

Rational ESG bondholders anticipate this moral hazard. Consequently, without further credit enhancements, bondholders will demand a massive ex-ante endogenous risk premium (expanding the CVA term $\lambda_1 D$), effectively pricing the equity holders out of the market and preventing the project from ever being funded. Therefore, the financial architecture requires a secondary mechanism not merely to stabilize cash flows, but to explicitly cure the agency conflict between founders and debt providers.

9.4 Risk Sharing via PPA Derivative Collars

To secure the debt tranche and prevent the SPV from facing insolvency during periods of low macroeconomic energy demand, the AI technology company must provide a credit enhancement. Instead of issuing a direct corporate debt guarantee, which would defeat the purpose of the asset-light SPV, the tech company structures the Power Purchase Agreement (PPA) with an embedded financial collar.

In this structure, the SPV sells its energy to the tech company at a floating market price P_t . However, the contract contains two crucial derivative-like provisions. First, the tech company contractually guarantees to purchase energy at no less than a specific floor price, P_{floor} . This floor is negotiated and set at a level high enough to ensure the SPV's revenues

will always be sufficient to cover its fixed debt coupon payments, thereby preventing default. In financial terms, the tech company has granted the SPV a perpetual American put option. In exchange for this valuable downside protection, the SPV agrees to cap the maximum energy price at P_{cap} , effectively granting an American call option back to the tech company.

The net present value of the SPV equity, $V_{SPV}(P_t)$, is therefore the standard project value plus the value of this embedded derivative collar:

$$V_{SPV}(P_t) = V_{standard}(P_t) + Put(P_t, P_{floor}) - Call(P_t, P_{cap}) \quad (52)$$

Because the energy demand is perpetual, the value of the put option is derived from its HJB equation. Let's break down its valuation. The value is split into two regimes depending on whether the option is exercised:

$$Put(P_t) = \begin{cases} \frac{A_{mono}^{osb} P_{floor}}{r_L} - \frac{A_{mono}^{osb} P_t}{r_L - \alpha} & \text{if } P_t \leq P_{floor} \\ B_1 P_t^{\beta_N} & \text{if } P_t > P_{floor} \end{cases} \quad (53)$$

The logic for these two cases is as follows:

- **Exercised Region ($P_t \leq P_{floor}$):** When the market price is at or below the floor, the SPV continuously "exercises" its right to sell at P_{floor} . The cash flow it receives at each instant is the difference between the guaranteed price and the market price: $A_{mono}^{osb}(P_{floor} - P_t)$. However, since the SPV is getting a fixed price, its effective value is no longer stochastic. The payoff is the present value of receiving a fixed annuity of $A_{mono}^{osb} P_{floor}$ forever (valued at $\frac{A_{mono}^{osb} P_{floor}}{r_L}$) minus the present value of the floating asset it must deliver, which is the standard project value $\frac{A_{mono}^{osb} P_t}{r_L - \alpha}$.
- **Unexercised Region ($P_t > P_{floor}$):** When the market price is above the floor, the option is not exercised. However, it still holds time value based on the possibility that the price could fall below the floor in the future. This "expectation value" must satisfy the HJB equation. As previously discussed, the general solution is $c_A P^\beta + c_B P^{\beta_N}$. For a put option, we need its value to approach zero as the underlying price P_t goes to infinity. The term P^β (with $\beta > 1$) would explode, so we must set its coefficient to zero. This leaves only the term $B_1 P_t^{\beta_N}$, where β_N is the negative root of the characteristic quadratic.

Why is β_N strictly negative? The characteristic quadratic equation from the HJB is $f(\beta) = \frac{1}{2}\sigma^2\beta(\beta-1) + \alpha\beta - r_L = 0$. This is a convex quadratic function. Since $f(0) = -r_L < 0$ and $f(1) = \alpha - r_L < 0$ (by the no-bubble assumption), the two real roots must lie on either

side of the interval $[0, 1]$. Therefore, the equation is guaranteed to have one positive root $\beta > 1$ and one negative root $\beta_N < 0$. The negative exponent ensures that as $P_t \rightarrow \infty$, the option value $B_1 P_t^{\beta_N} \rightarrow 0$, satisfying the necessary economic boundary condition.

To determine the unknown constant B_1 , I apply the classic smooth-pasting condition at the boundary P_{floor} . The smooth-pasting condition is a necessary first-order condition for the optimal exercise of an American-style option. It mandates that at the optimal exercise boundary, the derivative (the delta) of the option's value function must be equal to the derivative of its exercise payoff function. This ensures there is no kink at the boundary, which would imply a non-optimal exercise strategy. Taking the derivative of both cases with respect to P and setting them equal at P_{floor} yields:

$$B_1 \beta_N P_{floor}^{\beta_N - 1} = - \frac{A_{mono}^{osb}}{r_L - \alpha} \quad (54)$$

Isolating B_1 yields the exact valuation constant:

$$B_1 = - \left(\frac{A_{mono}^{osb}}{r_L - \alpha} \right) \frac{P_{floor}^{1 - \beta_N}}{\beta_N} \quad (55)$$

Because β_N is strictly negative and all other terms are positive, the constant B_1 is strictly positive. This is a crucial result: it means the put option has positive value and strictly increases the total value of the SPV equity. I mathematically prove that executing a PPA floor actively subsidizes the SPV, injecting synthetic capital directly into the firm's balance sheet. By absorbing the downside tail risk of terrestrial energy markets, the tech firm dramatically lowers the SPV's probability of coupon default, ensuring that institutional bondholders will confidently clear the debt tranche, effectively solving the \$14 billion space-based energy financing bottleneck without violating asset-light corporate mandates.

By capping the downside tail risk and flattening the volatility of the equity tranche's cash flows, the PPA derivative collar suppresses the founder's incentive to engage in risk-shifting behavior. It mathematically cures the agency friction, ensuring that institutional bondholders will confidently clear the debt tranche, effectively solving the space-based energy financing bottleneck without violating asset-light corporate mandates.

10 Two Factor Real Options: Stochastic Demand and Endogenous Launch Costs

The preceding sections utilize a one factor model where the capital expenditure I is deterministic and constant. In reality, the advent of reusable rocketry implies that launch costs are strictly decreasing over time subject to technological volatility. The decision to deploy an artificial intelligence center to space depends not merely on rising terrestrial energy demand, but also on the stochastic decline in payload costs. I extend the framework to a two factor real options model. Let the launch cost I_t follow a geometric Brownian motion with a negative drift μ_I and volatility σ_I :

$$dI_t = -\mu_I I_t dt + \sigma_I I_t dW_I \quad (56)$$

where the Wiener processes for energy demand and launch costs are correlated such that $\mathbb{E}[dW_P dW_I] = \rho dt$.

To determine the optimal value of the unexercised option $V(P_t, I_t)$, the firm must satisfy the two dimensional Hamilton-Jacobi-Bellman partial differential equation. Expanding the expectation operator utilizing the multi-dimensional Ito's Lemma yields:

$$\frac{1}{2}\sigma_P^2 P^2 V_{PP} + \frac{1}{2}\sigma_I^2 I^2 V_{II} + \rho\sigma_P\sigma_I P I V_{PI} + \alpha P V_P - \mu_I I V_I - rV = 0 \quad (57)$$

Because both the payoff and the strike price of the real option are homogenous of degree one in the state variables, the firm evaluates the ratio of the state variables. Let $Z_t = P_t/I_t$ represent the terms of trade between energy revenue and launch costs. Substituting the transformation $V(P, I) = I \cdot v(z)$ allows the extraction of the partial derivatives: $V_P = v'(z)$, $V_I = v(z) - z v'(z)$, $V_{PP} = \frac{1}{I} v''(z)$, $V_{II} = \frac{z^2}{I} v''(z)$, and $V_{PI} = -\frac{z}{I} v''(z)$. Substituting these derivatives back into the two dimensional partial differential equation and dividing the entire equation by I strictly reduces the model to a one dimensional ordinary differential equation:

$$\frac{1}{2}(\sigma_P^2 + \sigma_I^2 - 2\rho\sigma_P\sigma_I) z^2 v''(z) + (\alpha + \mu_I) z v'(z) - (r + \mu_I) v(z) = 0 \quad (58)$$

This mathematical reduction proves that the drift of the structural ratio is strictly positive, driven simultaneously by rising energy demand α and falling rocket costs μ_I . The firm faces a new characteristic root β_Z derived from the combined volatility profile. Applying the standard value matching and smooth pasting boundary conditions to the ratio variable Z_t , the optimal investment rule is no longer a fixed price scalar, but rather a linear free bound-

ary in the coordinate space. The firm invests the exact moment the energy price crosses the threshold boundary $P_t = z^* I_t$. This structural proof demonstrates that heavy engineering investments in rocket reusability, which increase the decay parameter μ_I , actively rotate the threshold boundary downward, drastically accelerating the economic viability of space based artificial intelligence data centers.

To ensure that the two-factor real option value does not contain speculative bubbles, the problem must strictly satisfy the transversality condition. As the temporal horizon approaches infinity, the discounted expected value of the asset must converge to zero. Let the discount factor under the equivalent martingale measure $\mathbb{Q}$ be defined by the risk-free rate r . The transversality condition requires:

$$\lim_{t \rightarrow \infty} \mathbb{E}^{\mathbb{Q}} [e^{-rt} V(P_t, I_t)] = 0 \quad (59)$$

Because the unexercised option value is bounded above by the active project value, it suffices to prove transversality for the active perpetuity. The expected value of the energy demand P_t grows at the risk-neutral drift $(\alpha - \lambda_{mkt}\sigma)$, where λ_{mkt} is the market price of risk, while the launch cost I_t decays at rate μ_I . Substituting the bounds into the expectation operator yields:

$$\lim_{t \rightarrow \infty} e^{-rt} \left(\frac{A_{mono} P_0 e^{(\alpha - \lambda_{mkt}\sigma)t}}{r - (\alpha - \lambda_{mkt}\sigma)} \right) = \lim_{t \rightarrow \infty} \left(\frac{A_{mono} P_0}{r - \alpha + \lambda_{mkt}\sigma} \right) e^{-(r - \alpha + \lambda_{mkt}\sigma)t} \quad (60)$$

For this limit to converge strictly to zero, the exponent must be negative, requiring the fundamental parametric inequality: $r > \alpha - \lambda_{mkt}\sigma$. Because the calibrated risk-free rate ($r = 0.0581$) strictly exceeds the risk-adjusted drift, the transversality condition holds. This mathematical proof guarantees that the two-factor boundary $P_t = z^* I_t$ represents a fundamentally sound economic equilibrium rather than a divergent mathematical artifact.

11 Empirical Calibration Strategy and Methodology

To bridge the gap between continuous-time theory and structural policy implementation, the mathematical parameters of this model must be rigorously calibrated using observable macroeconomic, financial, and aerospace datasets. This calibration strategy anchors the abstract stochastic jump-diffusion processes directly to real-world capital and astropolitical environments.

11.1 Estimating the Geometric Brownian Motion via MLE

The fundamental stochastic variable driving the investment timing game is the terrestrial demand for energy, modeled through the wholesale price P_t . The parameters governing this diffusion—specifically the drift α and the volatility σ —are calibrated using high-frequency wholesale electricity pricing data from the United States Energy Information Administration (EIA), mapping specifically to nodal pricing from the PJM Western Hub.

A primary empirical challenge in electricity markets is that standard hourly or daily spot prices are highly volatile and heavily dominated by transient, mean-reverting weather shocks (e.g., unexpected heatwaves). However, multi-billion dollar extraterrestrial infrastructure investments are not triggered by transient weather noise; they are evaluated against long-term, secular macroeconomic growth. Therefore, to correctly parameterize the continuous-time option value, I apply a 12-month trailing moving average filter to the raw EIA time series. This acts as a mathematical low-pass filter, stripping out short-term mean-reversion and perfectly isolating the secular Geometric Brownian Motion (GBM) expansion trend.

Given that the smoothed underlying price P_t evolves according to a GBM, the discrete log-returns over a monthly time step Δt , defined as $x_t = \ln(P_t/P_{t-\Delta t})$, are independent and normally distributed. Applying standard Maximum Likelihood Estimation (MLE), the discrete sample mean $\hat{\mu}$ and sample variance $\hat{s}^2$ of the log-returns are converted into continuous annualized parameters using Ito’s Lemma:

$$\hat{\sigma} = \frac{\hat{s}}{\sqrt{\Delta t}}, \quad \hat{\alpha} = \frac{\hat{\mu}}{\Delta t} + \frac{1}{2}\hat{\sigma}^2 \quad (61)$$

Executing this estimation on the deseasonalized PJM dataset yields an annualized volatility of $\sigma = 0.1443$ and an annualized drift of $\alpha = 0.0286$. A continuous structural drift of 2.86% accurately encapsulates the sustained demand shock imposed by data center load expansion and AI electrification.

11.2 Parameterizing Financial Constraints and Revenue Premiums

The financial parameters dictating the firm’s intertemporal discount factor are calibrated utilizing macroeconomic data from the Federal Reserve Economic Data (FRED) database. The standard baseline risk-free rate is set to $r = 0.0581$, reflecting the elevated, long-term 10-year Treasury yields required for extended, multi-decade capital horizons. The sustainable finance mechanism introduces an Asymmetric Greenium g , representing the yield concession offered by ESG-mandated institutional investors for sustainable asset deployment. Based

on the empirical yield spreads of frontier climate-tech subsidies and green municipal bonds, the greenium is parameterized at 50 basis points ($g = 0.0050$). This effectively reduces the cost of capital for the pioneer firm to an adjusted rate of $r_L = 0.0531$. Total baseline capital expenditure I is calibrated to \$14 Billion ($I = 100$ in normalized model units), utilizing cost-per-kilogram payload metrics published by commercial heavy-lift launch providers such as SpaceX to place a 2-million-kilogram SBSP constellation into operation.

Crucially, the cash-flow multiplier A must reflect the unique physical attributes of space-based energy. Terrestrial solar power is intermittent, suffering from the "duck curve" and yielding lower overall wholesale capture rates. Conversely, a Space-Based Solar Power constellation in geostationary orbit generates uninterrupted, 24/7 carbon-free baseload power. To empirically calibrate this advantage, I utilize Q4 2025 Power Purchase Agreement (PPA) pricing data from LevelTen Energy. Assuming a baseline wholesale average of \$50/MWh, terrestrial solar PPAs command \$61.7/MWh, while 24/7 baseload renewable assets (proxied by advanced wind/geothermal) command \$73.7/MWh.

Therefore, a monopolist SBSP pioneer capturing the 24/7 baseload premium is assigned a revenue multiplier of $A_{mono}^{osb} = \frac{73.7}{50} = 1.474$. If a second firm enters the same orbital sector and causes physical grid congestion, the premium is mathematically collapsed back to the standard intermittent solar proxy, yielding a duopoly multiplier of $A_{duo}^{osb} = \frac{61.7}{50} = 1.234$. This introduces intense endogenous price cannibalization into the model.

11.3 Estimating Endogenous Spatial Risk

Finally, the Poisson arrival rates of an orbital congestion collapse, λ_1 and λ_2 , parameterize the physical astropolitical tragedy of the commons. These structural jump probabilities are grounded using space situational awareness data and conjunction assessments from the European Space Agency (ESA) Space Debris Office. Let $y_{t,i}$ represent the count of critical orbital conjunctions (close approaches within a 1-kilometer radius) observed during spatial regime i . Assuming these conjunctions follow a non-homogeneous Poisson process, the event rate is scaled by $\phi = 10^{-3}$, a standard probabilistic risk assessment multiplier converting conjunction warnings into terminal Kessler Syndrome probabilities. This empirical framework establishes a baseline monopoly navigation risk of $\lambda_1 = 0.0050$. Upon the entry of a secondary firm, the inability to perfectly coordinate independent, proprietary navigation algorithms triggers an exponential spike in spatial density warnings, empirically shifting the duopoly collapse risk to $\lambda_2 = 0.0600$.

11.4 Calibrating Deep Space Computation and Reusability Parameters

To operationalize the spatial arbitrage and endogenous risk extensions, secondary parameters are calibrated using aerospace and energy agency projections. The growth rate of space based artificial intelligence computation α_D is parameterized at 0.1200, derived from the baseline compound annual growth rate of global data center electricity consumption published in the 2024 International Energy Agency electricity market report. The transmission efficiency of beaming power from deep space τ is set to 0.40, reflecting the inverse square law losses inherent in current microwave array demonstrations.

The capital expenditure penalty for deep space deployment is parameterized utilizing commercial lunar payload services contract data. Because escaping the gravity well requires substantially more delta-v propulsion, the payload cost per kilogram to Lagrange points is empirically observed to be multiple times higher than low Earth orbit. Consequently, I_{L1} is scaled to 2.5 times the baseline near Earth capital expenditure. Finally, the technological efficacy of in orbit servicing and debris capture κ is set to 0.05, proxying the marginal percentage reduction in structural failure probabilities per million dollars invested in autonomous navigation and robotic retrieval systems.

11.5 Calibrating Deep Space AI Demand and Payload Differentials

To calibrate the spatial arbitrage extension, the initial computation demand D_0 and its continuous drift α_D are estimated using commercial cloud compute pricing. Utilizing historical hourly lease rates for NVIDIA H100 tensor core GPUs from major cloud providers (e.g., Lambda Labs and AWS EC2), initial space-equivalent revenue is benchmarked against terrestrial equivalents. The drift parameter α_D is extracted from the Epoch AI “AI and Compute” historical dataset, which tracks the exponential growth in training floating-point operations per second (FLOPs). Applying a log-linear regression to this time series yields a conservative, risk-adjusted commercial demand drift of $\alpha_D = 0.1200$ (12% annually).

Furthermore, the capital expenditure penalty for deep space deployment (I_{L1}) is parameterized utilizing payload mass-to-orbit ratios published by commercial heavy-lift providers. Because escaping Earth’s gravity well to Lagrange points requires substantially more ΔV propulsion, a Falcon Heavy vehicle capable of delivering 63.8 metric tons to Low Earth Orbit can only deliver approximately 16 metric tons to deep space transfer orbits. This empirical payload ratio mathematically scales I_{L1} to roughly 2.5 times the baseline near-Earth capital

expenditure.

11.6 Parameterizing Reusability and Debris Capture Efficacy

The endogenous risk mitigation variables, specifically the capital allocation x and the technological efficacy κ , are anchored using aerospace public-private partnership data. The baseline investment x required to develop autonomous reusable and debris-capture systems is proxied using milestone funding data from NASA’s Commercial Orbital Transportation Services (COTS) program. To calibrate κ , which dictates the exponential decay of physical collision risk, I utilize the European Space Agency’s (ESA) Annual Space Environment Report. By regressing the historical reduction in required collision avoidance maneuvers against active satellite tracking upgrades, κ is conservatively parameterized at 0.05, implying a 5% compounding reduction in catastrophic failure probability per normalized unit of mitigation CapEx.

11.7 SPV Capital Structure and Two-Factor Launch Cost Drift

The financial architecture parameters are calibrated using macroeconomic debt indices and historical launch cost databases. The SPV senior debt rate r_D is anchored to the ICE BofA BBB US Corporate Index Effective Yield, accessed via the Federal Reserve Economic Data (FRED) database, adjusted downward by the empirical 50 basis-point greenium.

To calibrate the two-factor stochastic boundary, the negative drift of launch costs μ_I and its volatility σ_I are estimated using historical cost-per-kilogram data compiled by the Center for Strategic and International Studies (CSIS). The dataset tracks the historical transition from the Space Shuttle era ($\approx \$10,000/\text{kg}$) to the modern Falcon 9 era ($\approx \$1,500/\text{kg}$). Evaluating the discrete log-returns of this cost decay over a 15-year period yields a continuous negative drift of $\mu_I = 0.0400$ and a historical cost volatility of $\sigma_I = 0.1000$. The correlation coefficient $\rho = 0.20$ is assumed between energy demand and launch costs, reflecting the macroeconomic reality that booming technological demand slightly inflates aerospace supply chain costs.

12 Discussion of Numerical Results

Plugging the empirically calibrated vectors into the continuous-time real options mathematical engine yields a precise, dollar-value resolution to the mechanism design problem. By transitioning from abstract algebraic proofs to structural numerical estimation, the profound economic impact of the orbital-sustainability bond becomes distinctly quantifiable.

The outputs are presented comprehensively across three tables and three extensive graphical frameworks, each detailing a distinct structural consequence of asymmetric finance on the astropolitical commons.

12.1 Calibration and Equilibria Shifts

Table 1 explicitly logs the baseline input parameters derived through the empirical calibration strategy. The most critical integration in this matrix is the differentiation between the baseload PPA premium ($1.474x$) and the standard intermittent PPA premium ($1.234x$). By anchoring the continuous-time revenue functions to LevelTen Energy’s empirical spot pricing, the model transitions from a theoretical abstract into a highly actionable corporate finance valuation. Furthermore, identifying the baseline capital expenditure at \$14.00 Billion establishes the massive scale of the sunk costs (I) required to initiate extraterrestrial extraction. At this scale of irreversible capital deployment, even fractional deviations in optimal entry timing result in the loss of hundreds of millions of dollars, necessitating extreme precision in the formulation of the threshold equilibria.

Table 1: Empirical Parameter Calibration (LevelTen Q4 2025 Integration)

Parameter	Description	Value
α	Energy Demand Drift (EIA)	0.0286
σ	Energy Volatility (EIA)	0.1443
r_L	Greenium-Adjusted Rate (50 bps spread)	0.0531
A_{mono}^{osb}	Baseload PPA Premium (Wind Proxy $\frac{73.7}{50}$)	1.474x
A_{duo}^{osb}	Standard PPA Premium (Solar Proxy $\frac{61.7}{50}$)	1.234x
I	Baseline CapEx (SpaceX)	\$14.00B

Note: This table presents the empirically calibrated inputs. Demand parameters are derived via maximum likelihood estimation on PJM wholesale nodal pricing. Revenue multipliers reflect real-world Q4 2025 power purchase agreement spreads, capturing the exact baseload advantage of space-based solar power over terrestrial intermittent renewables.

Table 2 provides the comparative statics of the investment game, proving the fundamental mathematical existence of the pre-emptive race. In a theoretical vacuum where a firm holds a legally protected, unthreatened monopoly right, the optimal macroeconomic entry point is calculated at an energy price of $P_M^* = 8.37/\text{MWh}$. At this specific threshold, the marginal value of waiting perfectly balances the opportunity cost of foregone cash flows. However, because orbit represents an anarchic frontier entirely devoid of enforceable property rights, standard market competition forces the pioneer to pre-emptively deploy capital at a severely depressed price threshold of $P_L^* = 5.33/\text{MWh}$. The firm effectively cannibalizes its own upside to prevent a rival from seizing the market share.

By introducing the optimally calibrated bond, the model successfully engineers a profound double dividend. Subsidized by the 50-basis-point greenium ($r_L = 0.0531$) and the baseload revenue premium, the required hurdle rate of the pioneer drops aggressively from 5.33 down to $P_L^* = 2.98/\text{MWh}$. This achieves the first policy objective: accelerating the immediate deployment of clean energy infrastructure. Simultaneously, the heavy congestion penalty (C_p) inflates the required hurdle rate of the follower from 16.78 up to 21.81. Because the terrestrial energy drift (α) is constrained to a 2.86% continuous annual growth rate, this massive price gap extends the expected deterrence window—the expected timeframe a monopoly remains completely unchallenged—from a standard 40.1 years to a highly secure 69.6 years.

Table 2: Threshold Comparisons: Market Timing under Asymmetric Finance

Metric	Monopoly (Unthreatened)	Standard Race	OSB Optimal Mechanism
Leader Entry (P_L^*)	8.37	5.33	2.98
Follower Entry (P_F^*)	N/A	16.78	21.81
Deterrence Window (Years)	∞	40.1	69.6

Note: The deterrence window calculates the expected hitting time between the entry of the leader and the threshold of the follower based on the geometric drift $\alpha = 0.0286$. The mechanism successfully forces the follower out by nearly three decades, neutralizing the duopoly collision risk.

Table 3 quantifies the exact monetary value of recovering this economic inefficiency. In the standard, unconstrained race, the destruction of option value through premature entry drops the private ex-ante net present value of the firm to \$1.28 Billion. Under the bond market structure, the strategic delay of the follower firmly halts the pre-emptive race, allowing the leader to retain substantial, protected monopoly rents. This dynamic causes the private net present value to surge to \$1.46 Billion. Because the firm earns \$180 Million more under the restrictive environmental covenant than in the free market, adoption of the bond is mathematically proven to be a strictly dominant, subgame perfect strategy. Most importantly, the mechanism does not merely inflate private wealth; the expected present value of the contractual congestion liabilities reserves strictly \$0.41 Billion for the global equity fund. When combining the private NPV gain with the public equity trust, the total expected welfare expands from \$1.28 Billion to \$1.87 Billion, achieving an immense 21.7% recovery ratio of the total potential deadweight loss.

Table 3: Welfare Decomposition and Deadweight Loss Recovery

Metric (Billions USD)	Standard Market	OSB Market	Net Change
Private Firm NPV	\$1.28B	\$1.46B	+\$0.18B
Global Equity Fund	\$0.00B	\$0.41B	+\$0.41B
Total Welfare	\$1.28B	\$1.87B	+21.7% Recovery Ratio

Note: Ex-ante expected values are evaluated at an initial scarcity price of $P_0 = 3.0$. The recovery ratio quantifies exactly how much of the theoretical deadweight loss (the value destroyed by moving from a pure, unthreatened monopoly to a competitive race) is successfully clawed back by the mechanism.

12.2 Visualizing the Strategic Base Mechanism

Figure 1 visually corroborates the structural mechanics of the model across six highly interconnected dimensions. By translating the stochastic calculus into continuous state-space geometries, these panels definitively prove how asymmetric financial contracting governs competitive behavior in anarchic environments.

Panel A illustrates the absolute ex-ante welfare recovery evaluated at the baseline starting price of $P_0 = 3.0$. The solid red bar represents the heavily cannibalized firm value inherent in the standard pre-emptive race. The stacked blue and green bars of the optimal framework graphically prove the mechanism’s success in achieving practical Kaldor-Hicks efficiency. The visualization confirms that not only does the private firm gain net value, but a significant secondary surplus is generated purely for terrestrial social equity. This proves definitively that decentralized, market-driven financial instruments can perfectly substitute for sovereign taxation.

Panel B visualizes the exact mathematical proof of asymmetric rent retention. In a standard unregulated game, the value curve of the leader is forced intertemporally downward until it intersects the option curve of the follower, dissipating all first-mover advantage. However, because the bond grants a 50 basis point greenium and a baseload revenue multiplier exclusively to the certified pioneer, the value curve of the leader is shifted fundamentally higher. Consequently, the leader executes its option exactly where the option of the disadvantaged follower binds at $P = 2.98$. The significant vertical distance between the curves precisely at this threshold represents pure, protected economic rent that the pre-emptive race can no longer dissipate.




Figure 1: Structural Mechanics of the Orbital-Sustainability Bond across Six Dimensions.

Note: Panel A displays Kaldor-Hicks surplus recovery. Panel B provides the mathematical proof of halted pre-emption via asymmetric curves. Panel C demonstrates structural sensitivity to penalty calibration. Panel D simulates the stochastic timeline over a 40-year asset lifecycle. Panel E maps the Inverse Gaussian First Passage Time probability density. Panel F provides the mechanism designer's optimization heatmap for the competitive moat.

Panel C isolates the mathematical sensitivity of the Pigouvian entry delay mechanism. The x-axis tracks the structural magnitude of the congestion penalty (C_p), while the y-axis maps the resultant follower threshold (P_F^*). The strictly monotonic, linear upward slope proves the efficacy of the penalty lever. As the mechanism designer increases the contractual penalty for orbital degradation, the secondary firm is forced to wait for increasingly higher global energy prices to mathematically justify the risk of market entry. This structural geometry allows policymakers to precisely calibrate the penalty required to achieve any desired spatial deterrence window.

Panel D runs a rigorous 40-year Monte Carlo simulation of the energy transition timeline utilizing the empirically calibrated geometric Brownian motion parameters ($\alpha = 0.0286, \sigma = 0.1443$). Starting at $P_0 = 3.0$, the stochastic energy price (the jagged gray line) drifts upward subject to historical volatility. The leader establishes the space-based grid immediately at the blue dashed line (2.98). Over the next 40 years, despite sustained secular macroeconomic growth and highly volatile price peaks, the simulated state trajectory never touches the massive red deterrence barrier (21.81). The follower is actively, physically deterred from entering the orbit, ensuring the spatial state remains locked at $N_t = 1$, thus permanently averting the severe duopoly risk spike (λ_2).

Panel E elevates this analysis by mapping the first passage time probability density function. Utilizing the inverse Gaussian distribution inherent to the optimal stopping times of a geometric Brownian motion, this panel mathematically defines the risk horizon of a terminal Kessler Syndrome collapse. In the standard market (the red shaded area), the probability mass of the follower entering and triggering heavy congestion peaks critically around year 40. By implementing the bond (the green shaded area), the probability mass is structurally flattened and aggressively right-shifted. The statistical probability of triggering catastrophic orbital degradation within the standard operational lifespan of modern infrastructure is reduced to near-zero, successfully preserving the orbital commons for future generations.

Finally, Panel F provides a comprehensive heat map utilized by the mechanism designer to optimize the moat ratio (P_F^*/P_L^*). The x-axis scales the congestion penalty, while the y-axis scales the greenium spread. The dark purple regions in the bottom-left represent critically low moats resulting in heavy pre-emption, while the bright yellow regions in the top-right represent massive deterrence moats yielding near-absolute monopolies. This multi-dimensional matrix proves a critical economic reality: relying solely on a greenium (moving strictly up the y-axis) is insufficient to halt the race. A highly efficient, sustainable transition requires the simultaneous application of both carrots and sticks to push the market into the optimal, highly protected top-right quadrant.

12.3 Visualizing Spatial Arbitrage and Endogenous Mitigation

Figure 2 visualizes the structural dynamics of the spatial arbitrage framework and the endogenous reusability option, explicitly detailing how mechanisms alter the physical trajectory of the space economy.




Figure 2: Spatial Arbitrage, Future Cities, and Endogenous Risk Mitigation.

Note: Panel G identifies the spatial pivot threshold where the congestion cost of near-Earth orbit makes deep space deployment economically superior. Panel H simulates the growth option of the nascent extraterrestrial economy, modeling the demand driver (D_t) for future space urbanization. Panel I visualizes the U-shaped optimization of reusability and ISAM investment (x^*), where the firm balances direct CapEx against expected bond penalties. Panel J illustrates how endogenous mitigation allows for a sustainable reduction in the follower's entry barrier relative to static risk models.

Panel G proves the existence of the spatial pivot. The red line represents the net present value of entering the crowded near Earth orbit subject to the heavy penalty, while the purple line represents the net value of entering deep space. Due to the elevated capital expenditure ($I_{L1} = 250$), the deep space value is initially negative. However, precisely at the threshold where terrestrial energy price crosses $P = 36.8$, the heavy terrestrial penalty C_p mathematically forces the near Earth value curve beneath the deep space curve. At

this exact coordinate, the mechanism successfully repels the follower from near Earth orbit, inducing a Pareto-optimal deep space deployment of artificial intelligence data centers.

Panel H maps the foundational demand driver that makes deep space deployment viable: the extraterrestrial city growth option. Assuming a 12.0% continuous drift (α_D) derived from global data center power consumption rates, the model projects the exponential expansion of computational demand required to sustain lunar mining and orbital habitats. The shaded confidence interval represents the stochastic volatility inherent to frontier market expansion. This demonstrates that deep space deployment is not merely a penalty avoidance strategy, but a highly lucrative long-term real option on the future space economy.

Panel I provides the visual proof of the first-order condition for endogenous risk mitigation. The dashed gray line represents the linear, direct capital expenditure (x) invested in reusable rockets and servicing infrastructure. The dashed red line captures the expected present value of the penalty, which decays exponentially as the physical risk λ decreases. The solid dark line is the total expected friction. Because the mechanism sums a linear cost with a convex benefit, it guarantees a strict, U-shaped global minimum. The firm rationally invests exactly $x^* = 34.0$ to optimize its capital structure, proving that private firms will actively clean the orbit if the bond penalty is rigorously calibrated.

Finally, Panel J demonstrates how endogenous mitigation alters the timeline of the astropolitical transition. If Kessler Syndrome risk is assumed to be a static, exogenous constant, the follower is heavily delayed by the bond, requiring an entry threshold of $P_F = 21.81$. However, by utilizing debris capture technology to internalize the externality, the follower dramatically reduces its physical risk profile. This efficiency increases the present value of the firm, mathematically driving the required entry threshold down to $P_F = 12.10$. This proves that sustainable finance does not strictly halt progress; rather, it incentivizes technological innovations that lower the barriers to sustainable entry.

12.4 Visualizing Venture Capital Dilution and Launch Cost Dynamics

Figure 3 demonstrates the structural realities of financing space infrastructure utilizing equity dilution and two factor stochastic real options boundaries.

Panel K maps the endogenous dilution threshold $\eta^*(P_t)$. The red line represents standard venture capital commanding a 15% required rate of return. At low energy prices, the required dilution exceeds 100%, proving that premature market entry is mathematically unfinanceable. By introducing sovereign wealth co-investments that reduce the hurdle rate to 10% (the green line), the mechanism actively suppresses the required dilution, allowing founders




Figure 3: Venture Capital Contracting and Two-Factor Option Boundaries.

Note: Panel K maps the percentage of the firm founders must surrender to venture capital to fund capital expenditures. Panel L measures the retained equity value. Panel M illustrates the two factor real option boundary where firms evaluate both rising energy prices and falling rocket costs. Panel N simulates the stochastic decline of payload delivery costs via geometric Brownian motion.

to launch infrastructure earlier without surrendering controlling interest in the corporation.

Panel L confirms the maximization problem of the founders. The net present value of the retained equity scales linearly with the energy price. The vertical spread between the two lines perfectly captures the surplus value generated for the founders by securing the subsidized sovereign wealth capital structure.

Panel M provides the visual proof of the two factor partial differential equation. Unlike a standard real option which yields a fixed scalar price threshold, the integration of stochastic launch costs (I_t) yields a linear action boundary defined by $P_t = z^* I_t$. The firm occupies the blue region below the boundary, waiting optimally for either energy demand to rise or rocket prices to fall. The moment the coordinate pair crosses into the red region above the line, the firm exercises the option to launch the artificial intelligence center.

Finally, Panel N visualizes the empirical driver of the I_t state variable. Simulating a geometric Brownian motion with a negative continuous drift ($\mu_I = 0.0400$), the payload cost trajectory mirrors the historical cost reductions achieved by reusable booster technology. This panel confirms that engineering innovations on the launchpad act as a fundamental accelerator for the space based energy transition, independently driving the market toward the structural action boundary.

12.5 Visualizing Financial Architecture and Risk Sharing

Figure 4 vividly translates the mathematical derivations of compound real options, Special Purpose Vehicle (SPV) tranching, and derivative PPA collars into actionable corporate finance strategies.




Figure 4: Financial Architecture: Staged Real Options, SPV Tranching, and PPA Collars.

Note: Panel O demonstrates the reduction in the initial entry threshold achieved by utilizing a compound real options framework for a LEO demonstration phase. Panel P maps the required return on equity (r_E) against SPV debt leverage, incorporating the endogenous jump-to-default Credit Valuation Adjustment (CVA). Panel Q visualizes the synthetic capital injected into the SPV's balance sheet when a technology company provides an embedded American put option via a PPA price floor at $P = 2.5$.

Panel O confirms the success of the staged deployment strategy. The solid red bar represents the \$4.70/MWh threshold required if the firm attempts to deploy the full \$14 billion GEO infrastructure immediately. By structuring the transition as a compound real option, the firm separates the \$4.2 billion LEO demonstration phase from the terminal scale-up. Because the downside risk is isolated, the Phase 1 entry barrier mathematically collapses to \$3.46/MWh. This visualizes a critical policy reality: breaking megaprojects into sequential real options allows technology firms to initiate the space-based transition years earlier than standard static Net Present Value calculations would suggest.

Panel P visualizes the profound impact of the jump-to-default credit valuation adjustment on SPV capital structures. The dashed green line represents the base unlevered yield of

5.31%. However, as the SPV relies heavier on Green Debt (moving right along the x-axis), the fixed coupon obligations scale linearly. This leverage fundamentally alters the risk profile of the subordinated equity tranche. By the time debt leverage reaches 80%, the required return on equity (r_E) spikes exponentially past 7.0%. This curve perfectly graphs Modigliani-Miller Proposition II in an astropolitical context, proving that while debt tranching successfully offloads initial capital requirements to institutional bondholders, it demands severe equity yield premiums to compensate founders for the highly volatile residual cash flows.

Finally, Panel Q provides the visual proof of synthetic capital injection via PPA derivative collars. The dashed gray line represents the unhedged net present value of the SPV, which drops perilously toward zero as wholesale energy prices fall. The blue line represents the SPV value secured by a technology company's PPA floor guaranteed at $P = 2.5$. The green shaded region between them precisely maps the value of the embedded American put option ($B_1 P_t^{\beta_N}$). This visualization confirms that the PPA is not merely a revenue contract; it is a structural financial derivative. By capping the downside tail risk, the tech firm legally injects synthetic equity into the SPV, artificially elevating its valuation and preventing institutional bondholder default during macroeconomic demand troughs.

13 Conclusion

The structural transition to extraterrestrial energy generation is threatened by a dynamic, astropolitical market failure. In the absence of a global sovereign, symmetric pre-emptive competition for finite orbital and planetary resources leads inevitably to premature investment, extreme physical congestion, and the ultimate destruction of economic surplus. This paper proves that this tragedy of the commons can be successfully averted through decentralized, structural mechanism design.

By endogenizing the cost of capital and the environmental preferences of corporate technology buyers, I demonstrate that an orbital-sustainability bond serves as a subgame perfect commitment device. Firms strictly prefer to adopt this highly restrictive instrument over standard, unconstrained finance because the premium baseload revenues generated by constrained technology buyers financially overwhelm the expected cost of an orbital congestion penalty.

When rigorously calibrated using empirical wholesale energy demand drift from the EIA, aerospace CapEx metrics, and an observable 50 basis-point frontier greenium from the Federal Reserve, this mechanism yields a powerful double dividend. It successfully accelerates the deployment of essential clean energy infrastructure, depressing the initial entry threshold to \$2.98/MWh. Simultaneously, it successfully prices the endogenous spatial externality,

driving the secondary entrant's threshold out to \$21.81/MWh. As proven by inverse Gaussian first passage time analytics, this extreme Pigouvian delay ensures that secondary competitors do not enter during the primary infrastructure lifecycle, extending the deterrence window to nearly 70 years and effectively averting the exponential spike in Kessler Syndrome probabilities.

Beyond static entry delay, the expanded framework reveals two critical dynamic benefits of the mechanism. First, by properly calibrating the congestion penalty, the mechanism induces a spatial pivot. Secondary firms are mathematically incentivized to bypass crowded near-Earth orbits and instead deploy to deep space, acting as the foundational infrastructure for future extraterrestrial urbanization. Second, when firms are granted the agency to invest in reusable rockets and debris mitigation, the contract strictly incentivizes them to do so. The financial penalty forces the internalization of environmental costs, leading to an optimal private investment that actively reduces the global failure rate and sustainably lowers long-term entry barriers.

Furthermore, by integrating artificial intelligence computation directly into the deep space arrays, the model demonstrates that firms can bypass the physical transmission limits of energy beaming. By processing data locally in the zero cost thermal environment of deep space and beaming only the resultant data back to Earth, human civilization can structurally decouple computational advancement from terrestrial climate degradation.

By expanding the theoretical framework into structural corporate finance, this research provides a comprehensive blueprint for funding frontier megaprojects. I demonstrate that breaking the massive \$14 billion infrastructure requirement into sequential LEO-to-GEO compound real options drastically lowers the initial barrier to entry, accelerating the timeline for early technological demonstration. Furthermore, by structuring the capital within a Special Purpose Vehicle (SPV), the model proves that Power Purchase Agreements with embedded price floors act as synthetic equity injections. These derivative collars mathematically cure the asset substitution problem, protecting institutional bondholders from default risk and enabling the secure tranching of sustainable debt alongside two-factor stochastic launch cost dynamics.

This framework fundamentally halts the value-destroying pre-emptive race. It expands total expected welfare from \$1.28 billion to \$1.87 billion, achieving an immense 21.7% recovery of potential deadweight loss. Furthermore, the mechanism contractually transfers \$410 million of this recovered surplus directly into a global equity fund for terrestrial climate adaptation. Ultimately, this research proves that sophisticated financial contracting can seamlessly align the pursuit of private, astropolitical technological expansion with the preservation of the extraterrestrial commons and the equitable distribution of global re-

sources.

References

- Adilov, N., Alexander, P. J., and Cunningham, B. M. (2015). An economic analysis of earth orbit pollution. *Environmental and Resource Economics*, 60(1):81–98.
- Bolton, P. and Kacperczyk, M. (2021). Do investors care about carbon risk? *Journal of Financial Economics*, 142(2):517–549.
- Dixit, A. K. and Pindyck, R. S. (1994). *Investment Under Uncertainty*. Princeton University Press.
- Flammer, C. (2021). Corporate green bonds. *Journal of Financial Economics*, 142(2):499–516.
- Grenadier, S. R. (2002). Option exercise games: An application to the equilibrium investment strategies of firms. *The Review of Financial Studies*, 15(3):691–721.
- Huisman, K. J. (2013). *Technology investment: A game theoretic real options approach*, volume 28. Springer Science & Business Media.
- Meckling, W. H. and Jensen, M. C. (1976). Theory of the firm. *Managerial behavior, agency costs and ownership structure*, 3(4):305–360.
- Nordhaus, W. D. (2014). Estimates of the social cost of carbon: concepts and results from the dice-2013r model and alternative approaches. *Journal of the Association of Environmental and Resource Economists*, 1(1/2):273–312.
- Ostrom, E. (1990). *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press.
- Pástor, Ľ., Stambaugh, R. F., and Taylor, L. A. (2021). Sustainable investing in equilibrium. *Journal of Financial Economics*, 142(2):550–571.
- Rao, A., Burgess, M. G., and Kaffine, D. (2020). Orbital-use fees could more than quadruple the value of the space industry. *Proceedings of the National Academy of Sciences*, 117(23):12756–12762.
- Smit, H. T. and Trigeorgis, L. (2012). *Strategic investment: Real options and games*. Princeton University Press.
- Weitzman, M. L. (2009). On modeling and interpreting the economics of catastrophic climate change. *The review of economics and statistics*, 91(1):1–19.