

Financing the Space-Based Energy Transition: Endogenous Risk Mitigation, Spatial Arbitrage, and Sustainable Mechanism Design

Shiming Wu

Abstract

The computational requirements of modern artificial intelligence models and data centers are driving an unprecedented transition toward extraterrestrial energy generation, specifically space-based solar power. However, prime near-Earth orbits and lunar sectors represent finite common-pool resources subject to endogenous congestion risk based on total spatial occupancy. This paper develops a continuous-time stochastic differential game to analyze the extraction of these frontier resources. I prove that standard market competition yields a pre-emptive investment race that destroys substantial economic surplus and collapses into a Pareto-inefficient tragedy of the commons. Because extraterrestrial jurisdictions lack a central enforcing sovereign, traditional Pigouvian taxes are unimplementable. Consequently, I design a decentralized mechanism—the orbital-sustainability bond—that leverages an asymmetric greenium and corporate power purchase agreement premiums to endogenize spatial externalities. I derive the closed-form Hamilton-Jacobi-Bellman equations and rigorously calibrate the model. I prove that an optimally calibrated bond acts as a subgame perfect commitment device, generating a double dividend: it accelerates the initial deployment of clean energy infrastructure by depressing the entry threshold to $P = 2.98$, while simultaneously delaying the entry of secondary competitors to $P = 21.81$. Furthermore, the model is extended to incorporate spatial arbitrage and endogenous risk mitigation. I demonstrate that the bond penalty induces a spatial pivot, incentivizing secondary entrants to deploy artificial intelligence data centers directly to deep space to utilize vacuum cooling and supply nascent extraterrestrial urbanization rather than crowding near-Earth orbits. Additionally, when physical congestion risk is made endogenous, the mechanism forces private firms to optimally invest in reusable rockets and debris capture technology, actively lowering the global failure rate. Finally, I solve a two-factor partial differential equation integrating stochastic payload costs alongside venture capital contracting to establish the optimal endogenous equity capital structure. This

framework halts the value-destroying pre-emptive race, yields a 21.7% recovery ratio of deadweight loss, and contractually allocates over \$400 million to a global terrestrial equity fund.

1 Introduction

The advancement of frontier technology is fundamentally bounded by physical energy capacity. The computational demands of training and operating large computing models have created a severe terrestrial energy bottleneck. A single next-generation chip consumes approximately 64.8 kilowatt-hours daily, and with technology firms planning thousands of new facilities, digital infrastructure is projected to account for a significant double-digit percentage of all United States electricity usage by the end of the decade. This demand shock is colliding with a rigid terrestrial grid, prompting sovereign nations and multinational corporations to direct capital toward space-based solar power and extraterrestrial energy generation, which promises uninterrupted, baseload renewable power.

However, this transition introduces a fundamental problem of industrial organization, mechanism design, and astropolitics. The optimal near-Earth orbits and strategic planetary locations represent finite common-pool resources. A critical aspect of these resources is that congestion is strictly endogenous to total spatial occupancy. A monopolist scaling infrastructure to meet consumer demand inherently increases its own risk of triggering an orbital cascade. Furthermore, when a second firm enters the same orbital sector, the difficulties of inter-firm operational coordination cause the risk of a terminal Kessler Syndrome collapse to spike exponentially. Kessler Syndrome refers to a theoretical scenario where the density of objects in low Earth orbit is high enough that collisions between objects cause an irreversible kinetic cascade, rendering specific orbital ranges inaccessible for generations.

The ultimate objective of this spatial transition extends beyond merely beaming power to terrestrial grids. As artificial intelligence models scale exponentially, they face absolute physical limits on Earth regarding thermal dissipation and land utilization. Consequently, this paper posits that the optimal long-term equilibrium involves migrating the data centers themselves into the extraterrestrial commons. By co-locating artificial intelligence computation centers directly with space-based solar arrays, firms can utilize the near absolute zero temperature of the space vacuum for costless thermal management. This paradigm shift permanently alleviates terrestrial global warming vectors and guarantees long-term energy safety by entirely removing industrial computational loads from the fragile Earth biosphere. While deploying computation to deep space currently exhibits higher capital expenditures than utilizing terrestrial deserts or nuclear facilities, the extraterrestrial paradigm represents

an infinite growth horizon that mathematically surpasses the rigid physical boundaries of the Earth.

A natural regulatory response to common-pool resource degradation in environmental economics is a centralized Pigouvian tax or a cap-and-trade quota system. However, extraterrestrial domains present a unique jurisdictional challenge. Governed loosely by the 1967 Outer Space Treaty, no single sovereign possesses the absolute authority to enforce mandates, deny access, or levy taxes in orbital space. Preventing the physical degradation of these resources therefore requires a fully decentralized, market-driven mechanism. Consequently, this paper structures this challenge around six core economic research questions. First, in a continuous-time stochastic game where collapse risk depends non-linearly on spatial occupancy, why does standard market competition inevitably lead to a Pareto-inefficient pre-emptive race? Second, how can sustainable finance mechanisms, acting as private commitment devices, internalize these spatial externalities without state intervention? Third, does the introduction of sustainable finance always yield optimal results, or can poorly calibrated green contracts paradoxically accelerate the tragedy of the commons? Fourth, under what precise empirical conditions does an asymmetric greenium generate a double dividend that aligns private technological expansion with the preservation of the astropolitical commons? Fifth, how does the introduction of deep space deployment alternatives transform the pre-emptive race into a problem of spatial arbitrage? Finally, can financial contracts mathematically induce private capital investment into reusable rockets and debris capture technology to actively suppress the physical failure rate?

This preservation aims to restore Kaldor-Hicks efficiency, a criterion utilized to evaluate whether a change in resource allocation improves aggregate social welfare. Under standard Kaldor-Hicks theory, a transition is considered efficient if the aggregate gains to the winners are large enough that they could, in theory, compensate the losers, resulting in a net social improvement. While traditional macroeconomic models often leave this compensation as a theoretical footnote, this paper proposes a structural financial mechanism that transforms this potential into a tangible transfer of wealth. By internalizing the spatial externality, the mechanism recovers immense deadweight loss and contractually captures a defined portion of that surplus—projected in this calibration at \$410 million—for a global equity fund aimed at terrestrial climate adaptation.

Regarding methodology, I develop a continuous-time real options game incorporating Poisson jump-diffusion processes to model state-dependent spatial collapse. The parameters of the stochastic process are explicitly calibrated using maximum likelihood estimation on wholesale electricity data from the EIA and PJM Interconnection. Furthermore, I integrate empirical power purchase agreement pricing from LevelTen Energy to differentiate

intermittent solar pricing from baseload pricing. Endogenous orbital failure probabilities are parameterized using European Space Agency conjunction assessments. I solve the Hamilton-Jacobi-Bellman equations to identify the value of the project across three distinct regimes: monopoly scarcity, duopoly abundance, and terminal collapse. The model is subsequently extended to incorporate mutually exclusive spatial options driven by a secondary geometric Brownian motion representing future extraterrestrial city demand, alongside a static optimization problem where firms dynamically select capital allocations for debris mitigation. Furthermore, to address the reality that space infrastructure requires external funding, I construct a structural model of venture capital contracting to establish endogenous equity dilution thresholds. Finally, I utilize a two-factor partial differential equation framework to integrate the stochastic decline of rocket launch costs alongside rising energy demand.

The quantitative results of this paper demonstrate that an unconstrained market yields a symmetric pre-emptive race where firms invest prematurely. Specifically, while a theoretical unthreatened monopolist would wait for an energy price threshold of \$8.37/MWh, the threat of competition forces initial entry at a depressed threshold of \$5.33/MWh, severely cannibalizing expected surplus. I find that the orbital-sustainability bond completely resolves this through asymmetric finance. Protected by an empirically grounded 50 basis-point greenium and a 1.474x baseload revenue premium, the first-mover accelerates clean energy deployment, dropping the entry threshold to \$2.98/MWh. Simultaneously, the heavy contractual penalty imposes a Pigouvian delay on the secondary entrant, driving the follower’s threshold from \$16.78 out to \$21.81/MWh.

The extended framework reveals that this heavy terrestrial penalty induces a spatial pivot, mathematically redirecting the secondary firm to deploy deep space artificial intelligence centers to capture the long-term growth option of extraterrestrial urbanization. Furthermore, by making physical risk endogenous to capital expenditure, I prove the mechanism forces the firm to invest an optimal \$34.0 million equivalent into reusable systems, which drops the physical failure rate and sustainably lowers the long-term entry barrier back to \$12.10/MWh. Through rigorous mathematical derivations of the capital structure, I demonstrate that sovereign wealth co-investments can substantially reduce required founder dilution, preventing premature funding failures. Ultimately, stochastic simulations prove that this comprehensive mechanism secures the orbital commons and yields a 21.7% total welfare recovery ratio.

The remainder of this paper proceeds as follows. Section 2 reviews the literature on real options, climate economics, and sustainable finance. Section 3 constructs the mathematical environment and derives the baseline pre-emptive race. Section 4 introduces the orbital-sustainability bond, detailing the subgame perfect proofs for asymmetric adoption. Sections

5 and 6 discuss spatial arbitrage and endogenous debris mitigation. Section 7 derives the venture capital equity contracting model. Section 8 solves the two-factor stochastic boundary problem. Section 9 outlines the empirical calibration strategy. Section 10 provides a discussion of the numerical and visual results. Section 11 concludes.

2 Literature Review

This paper first builds upon and extends the real options literature concerning strategic investment under uncertainty. Foundational work by [Dixit and Pindyck \(1994\)](#) established the mathematical apparatus for evaluating irreversible investments, which was subsequently extended into game-theoretic duopoly settings by [Huisman \(2013\)](#) and [Grenadier \(2002\)](#) to analyze pre-emptive investment races. A critical gap in this existing literature is that models of pre-emptive races typically focus strictly on market share competition and price cannibalization, assuming the underlying physical asset remains intact. This literature generally ignores the physical degradation of shared common-pool resources. My paper contributes to this field by integrating a permanent, state-dependent jump-risk collapse directly into the continuous-time pre-emption framework, demonstrating how total spatial occupancy fundamentally alters stochastic hitting times and optimal option exercise thresholds.

Second, this research intersects with the broad economic literature on climate change and the governance of common-pool resources. [Ostrom \(1990\)](#) demonstrated how institutions can govern the commons, while modern environmental macroeconomics led by [Weitzman \(2009\)](#) and [Nordhaus \(2014\)](#) focuses on pricing catastrophic tipping points. The primary limitation of these macroeconomic frameworks is their reliance on the existence of a benevolent social planner, a sovereign capable of enforcing Pigouvian taxes, or unified global climate treaties. These assumptions fail completely in anarchic international domains such as orbital space and extraterrestrial extraction. I contribute to this literature by proposing a fully decentralized, market-driven mechanism that achieves endogenous entry deterrence without a central planner, proving that private capital markets can substitute for state enforcement in geopolitical frontiers.

Finally, this paper contributes to the rapidly growing field of sustainable finance and environmental, social, and governance contracting. Recent theoretical work by [Pástor et al. \(2021\)](#) has formalized how investor preferences generate a greenium and lower the cost of capital for sustainable firms. However, a significant gap remains in structural corporate finance. Much of the literature focuses on the empirical asset pricing of the greenium or addresses moral hazard through principal-agent models, lacking a dynamic framework that shows how sustainability covenants actively alter strategic investment timing among oligopolistic com-

petitors. I contribute by providing a subgame perfect microfoundation that proves firms will voluntarily accept restrictive covenants not out of altruism, but as a strategic weapon to manipulate competitor entry thresholds and extend their own monopoly rents.

3 Model Environment and the Standard Equilibrium

I model the investment decision as a dynamic game within a three-regime market lifecycle. The state of the market depends on the number of active firms, denoted as $N_t \in \{0, 1, 2\}$, and a potential terminal congestion shock. The stochastic terrestrial energy price P_t , representing the demand for computation, follows a geometric Brownian motion:

$$dP_t = \alpha P_t dt + \sigma P_t dW_t \quad (1)$$

where α is the drift parameter representing secular growth in demand, σ is the volatility parameter representing uncertainty, and dW_t is the increment of a standard Wiener process. The constant risk-free interest rate is r , and I assume $r > \alpha$ to ensure the present value of future cash flows converges to a finite number. Firms face a sunk capital cost I to deploy extraterrestrial energy infrastructure.

3.1 The Three-Regime Market Lifecycle and Endogenous Risk

Unlike traditional capacity expansion models where exogenous shocks arrive uniformly, the risk of a terminal congestion collapse in this model is strictly endogenous to total spatial occupancy. The market transitions sequentially through three distinct physical regimes.

Regime 1 is the scarcity phase where $N_t = 1$. The first firm, defined as the Leader, enters the market. Operating without spatial competition, it captures a monopoly cash flow multiplier of $A_{mono}^{std} P_t$. Because a single firm operates with largely unrestricted orbital slots, it can optimize its own navigation and positioning algorithms without external interference. Therefore, it generates only a minimal baseline risk of a self-induced orbital cascade. I model this risk as a Poisson process with a very low arrival rate λ_1 .

Regime 2 is the abundance phase where $N_t = 2$. The second firm, defined as the Follower, enters the market. Competition for prime orbital real estate, alongside standard energy market share competition, drives down profit margins. This yields a lower cash flow multiplier $A_{duo}^{std} P_t$ for both active firms, where $A_{duo}^{std} < A_{mono}^{std}$. Crucially, the total spatial occupancy doubles. The physical crowding and the inability to perfectly coordinate independent, proprietary navigation algorithms cause the risk of a terminal Kessler Syndrome collapse to spike exponentially. The Poisson arrival rate jumps to λ_2 , where $\lambda_2 \gg \lambda_1$. This jump in

risk represents the physical manifestation of the tragedy of the commons, where individual rational entry leads to collective physical ruin.

Regime 3 is the congestion collapse. The orbital environment becomes physically unstable due to a Kessler Syndrome event. Upon the arrival of the shock, which occurs at rate λ_1 during the monopoly phase or λ_2 during the duopoly phase, operational efficiency is permanently impaired. The revenue parameter collapses to A_{coll}^{std} , which is strictly less than A_{duo}^{std} .

To establish the baseline of the model, I first derive the value of an active firm strictly in the terminal collapse regime. Because there are no further state transitions after a collapse, the value of the firm is simply the expected present value of a perpetual cash flow growing at rate α . Applying the standard Gordon Growth continuous-time integral for a geometric Brownian motion yields the collapsed project value:

$$V_{coll}^{std}(P_t) = \int_0^\infty \mathbb{E}[A_{coll}^{std} P_{t+\tau}] e^{-r\tau} d\tau = \frac{A_{coll}^{std} P_t}{r - \alpha} \quad (2)$$

3.2 Derivation of Risk-Adjusted Regime Values

Consider the market where firms utilize standard finance without any sustainability covenants. To find the value of a firm operating prior to collapse, I must establish intertemporal arbitrage conditions for both the scarcity and abundance regimes.

By the no-arbitrage principle, the required return on the asset over a differential time step dt must exactly equal the cash flow generated plus the expected capital gain. Let $\lambda_i \in \{\lambda_1, \lambda_2\}$ be the state-dependent risk. The Hamilton-Jacobi-Bellman equation is:

$$rV_{regime}^{std}(P_t)dt = A_{regime}^{std}P_tdt + \mathbb{E}[dV_{regime}^{std}(P_t)] \quad (3)$$

Because the asset is subject to both continuous price diffusion and discrete collapse jumps, I apply Ito's Lemma for jump-diffusion processes to expand the expected change term. The continuous geometric Brownian motion component is expanded using a second-order Taylor series, yielding $\alpha P_t V' dt + \frac{1}{2} \sigma^2 P_t^2 V'' dt$. The discrete jump component represents the probability $\lambda_i dt$ that the firm suffers a discrete capital loss, falling from the active regime value to the collapsed state. This capital loss is modeled as $\lambda_i (V_{coll}^{std} - V_{regime}^{std}) dt$.

Substituting these components back into the Hamilton-Jacobi-Bellman equation and dividing by dt yields the fundamental non-homogeneous linear differential equation:

$$\frac{1}{2} \sigma^2 P_t^2 (V_{regime}^{std})'' + \alpha P_t (V_{regime}^{std})' - (r + \lambda_i) V_{regime}^{std} + A_{regime}^{std} P_t + \lambda_i V_{coll}^{std} = 0 \quad (4)$$

To solve this differential equation, I guess a particular solution of the form $V = kP_t$ because the underlying cash flows are strictly linear in P_t . Taking the first derivative $V' = k$ and the second derivative $V'' = 0$, I substitute these into the equation. Replacing V_{coll}^{std} with the expression from Equation 2 and solving for the multiplier k yields the fundamental, risk-adjusted present value of the active firm in any regime subject to spatial risk λ_i :

$$V_{risky}^{std}(P_t, \lambda_i, A_{regime}^{std}) = \frac{A_{regime}^{std} P_t}{r + \lambda_i - \alpha} + \frac{\lambda_i}{r + \lambda_i - \alpha} \left(\frac{A_{coll}^{std} P_t}{r - \alpha} \right) \quad (5)$$

This derivation explicitly proves the economic intuition of the model. Higher spatial occupancy, represented by the shift from λ_1 to λ_2 , directly cannibalizes the present value of the firm by exponentially increasing the discount factor in the denominator of the primary cash flows.

3.3 The Follower's Option and Standard Entry

Before entering the market, the Follower holds a perpetual American call option to invest. Because the option itself yields no cash flows prior to exercise, its value $F(P_t)$ satisfies the standard homogeneous differential equation $0.5\sigma^2 P_t^2 F'' + \alpha P_t F' - rF = 0$. Ruling out asset bubbles and noting that the option value must approach zero as the energy price approaches zero, the standard solution is $F(P_t) = cP_t^\beta$, where $\beta > 1$ is the positive root of the characteristic quadratic equation $0.5\sigma^2 \beta(\beta - 1) + \alpha\beta - r = 0$.

The Follower optimizes its entry timing by selecting an optimal threshold P_F^* . Upon entry, the market shifts to the abundance phase, triggering the higher spatial risk λ_2 . At the optimal threshold, two economic boundary conditions must hold. First, value-matching requires that the option value equals the net payoff of investment: $F(P_F^*) = V_{risky}^{std}(P_F^*, \lambda_2, A_{duo}^{std}) - I$. Second, smooth-pasting ensures marginal optimization, meaning the derivative of the option equals the derivative of the payoff: $F'(P_F^*) = (V_{risky}^{std})'(P_F^*, \lambda_2, A_{duo}^{std})$.

Substituting the option form cP^β into the smooth-pasting equation allows me to solve for the constant c . Substituting this c back into the value-matching equation allows for the isolation of the threshold variable. Solving algebraically yields the Follower's standard entry threshold:

$$P_F^* = \frac{\beta}{\beta - 1} \frac{I}{(V_{risky}^{std})'(1, \lambda_2, A_{duo}^{std})} \quad (6)$$

3.4 The Leader's Pre-Emptive Race and Value Destruction

A firm contemplating entering the market as the Leader must evaluate the continuous threat of pre-emption. Its post-investment project value, $V_L(P_t)$, equals the value of a perpetual monopoly operating under the low risk λ_1 , minus the expected present value of the competitive loss triggered at the first hitting time when the price reaches the Follower's threshold. Utilizing the expected stochastic discount factor $(P_t/P_F^*)^\beta$ for a geometric Brownian motion reaching an upper boundary, the Leader's value is:

$$V_L(P_t) = V_{risky}^{std}(P_t, \lambda_1, A_{mono}^{std}) - \left(\frac{P_t}{P_F^*}\right)^\beta [V_{risky}^{std}(P_F^*, \lambda_1, A_{mono}^{std}) - V_{risky}^{std}(P_F^*, \lambda_2, A_{duo}^{std})] \quad (7)$$

The net present value of leading is defined as $L(P_t) = V_L(P_t) - I$. In a theoretical market where a firm holds an exogenous, unthreatened monopoly right, the firm would maximize its value by finding the absolute peak of the firm value curve. This point is called the socially optimal unthreatened threshold, representing the exact point at which a single firm, facing no threat of pre-emption or competitive entry, would optimally choose to invest. At this threshold, the marginal value of waiting for higher energy prices is exactly zero, maximizing the absolute net present value of the asset.

However, in standard market conditions, firms are symmetric and lack exclusive extraction rights. Because each firm holds an identical option to lead and fears being relegated to the less profitable Follower role, they engage in a symmetric pre-emptive race. They will continually deploy capital earlier to undercut their rival. Mathematically, this competitive race terminates strictly at the equilibrium threshold P_L^* where the economic rent of being the Leader is completely dissipated down to the option value of following. This is defined by the absolute Value-Matching condition:

$$L(P_L^*) = F(P_L^*) \implies V_L(P_L^*) - I = [V_{risky}^{std}(P_F^*, \lambda_2, A_{duo}^{std}) - I] \left(\frac{P_L^*}{P_F^*}\right)^\beta \quad (8)$$

Because the option to follow $F(P)$ is strictly positive and monotonically increasing, setting $L(P_L^*) = F(P_L^*)$ forces the intersection to occur strictly on the upward-sloping portion of the Leader's value curve. This mathematical property guarantees that the pre-emptive threshold is significantly lower than the monopoly optimum ($P_L^* \ll P_M^*$). The standard market yields a Pareto-inefficient race, deploying capital prematurely when market demand is too low to justify the expenditure, severely destroying expected total surplus in a tragedy of the commons.

4 Mechanism Design: Greenwashing vs. Pigouvian Delay

To correct this market failure without a central sovereign, I design the orbital-sustainability bond. This mechanism utilizes two distinct levers to alter the strategic landscape: a clean energy premium and an asymmetric greenium.

A greenium, or green premium, is the reduction in the cost of debt or cost of capital that institutional investors offer to firms that commit to strict, verifiable environmental covenants. Because sustainable funds possess mandates to allocate capital toward climate solutions, they accept slightly lower financial yields in exchange for quantifiable environmental impact, effectively lowering the upfront capital expenditure for the issuing firm.

If a firm issues the bond, verified spatial sustainability allows the firm to sign premium power purchase agreements with consumers, yielding elevated revenue parameters $A^{osb} > A^{std}$. A crucial condition is embedded in the contract: if an orbital congestion collapse occurs, the sustainability certification is immediately revoked. Consequently, the collapsed revenue parameter reverts to the standard baseline ($A_{coll}^{osb} = A_{coll}^{std}$). Consumers will not pay a premium for energy generated from a degraded orbit. In exchange for the upfront financial benefits, both firms bind themselves to a heavy congestion penalty C_p , which is diverted to a global equity fund upon collapse.

The global equity fund is a theoretical contractual trust designed to capture the congestion penalties levied upon firms that trigger orbital degradation. Instead of allowing the surplus to be lost to deadweight inefficiency, this fund contractually redistributes the capital to terrestrial climate adaptation projects, ensuring an equitable distribution of the extraterrestrial commons. Unlike purely theoretical welfare measures, the global equity fund represents an actual physical transfer of capital from the extractive firms to global society. This ensures that the recovery of deadweight loss from the pre-emptive race is shared globally, thereby fulfilling the Kaldor-Hicks efficiency criterion by transforming a potential social gain into a tangible social fund.

Because the penalty C_p is a fixed nominal liability paid at the random future time of collapse, its expected present value at the time of entry requires integrating the exponential probability density function against the continuous discount rate. The expectation for a firm operating in regime i with risk λ_i is:

$$\mathbb{E}[PV(C_{pi})] = \int_0^\infty e^{-rt}(1 - \theta)C_p\lambda_i e^{-\lambda_i t} dt = \frac{\lambda_i}{r + \lambda_i}(1 - \theta)C_p \quad (9)$$

4.1 The Mathematical Definition of Greenwashing

The introduction of sustainable finance does not automatically resolve the tragedy of the commons. If the contract is poorly calibrated, it creates a dual failure. Both the Leader and the Follower are capable of issuing bonds. Greenwashing is defined mathematically as a state where the financial benefits of the bond, specifically the clean energy premiums and capital reductions, strictly exceed the expected environmental penalty for the secondary entrant.

Under these parameters, the bond acts as a financial subsidy that artificially inflates the asset's value without deterring risky behavior. This causes the denominator of the Follower's entry threshold equation to expand faster than the numerator. Consequently, the Follower is prompted to enter at a lower price threshold than in the standard unregulated market ($\tilde{P}_F^* < P_F^*$). By incentivizing earlier entry, the poorly calibrated green bond accelerates the transition to the high-risk duopoly phase λ_2 , directly accelerating the physical degradation of the shared orbital resource.

4.2 The Asymmetric Greenium and Pigouvian Delay

To prevent greenwashing and restore efficiency, the optimally designed mechanism relies on a fundamental asymmetry in sustainable finance. ESG institutional investors rationally allocate greenium capital exclusively to the pioneer who establishes the essential clean infrastructure, represented by the Leader, such that $g_L > 0$. They strictly deny the greenium to the Follower ($g_F = 0$) because the secondary entry provides redundant capacity while actively spiking the spatial collapse risk.

The Follower, facing no greenium, evaluates its new entry threshold under the bond covenant. By calibrating C_p to be massive, the mechanism designer inflates the Follower's fixed cost barrier. The numerator in the threshold equation, which includes the sunk cost and the expected penalty, strictly dominates the enhanced revenue multiplier in the denominator. This elevates the Follower's hurdle rate, imposing a Pigouvian delay ($\tilde{P}_F^* > P_F^*$). The mechanism ensures the orbit is only crowded when market demand is extraordinarily high.

A critical theoretical hurdle is explaining why a cost-disadvantaged Follower would voluntarily adopt a high-penalty bond. The Follower chooses the bond framework over standard finance because consumers offer a significant price premium A_{duo}^{osb} for certified energy. At the moment of entry $\tilde{P}_F^*$, the perpetuity value of this premium completely overwhelms the one-time fixed penalty:

$$V_{risky}^{osb}(\tilde{P}_F^*, \lambda_2) - I - \mathbb{E}[PV(C_{p2})] > V_{risky}^{std}(\tilde{P}_F^*, \lambda_2) - I \quad (10)$$

Because this condition holds strictly, the Follower rationally restricts itself to the mechanism.

4.3 Proof of Strategic Dominance and the Double Dividend

The Leader strictly prefers the optimal bond because it acts as a strategic moat. Because the solo-risk λ_1 is negligible, the expected penalty poses almost zero upfront cost to the Leader ($\mathbb{E}[PV(C_{p1})] \approx 0$).

Crucially, the asymmetric greenium mathematically halts the destructive pre-emptive race. Let Firm B be the disadvantaged Follower ($g_F = 0$). Firm B will only preempt the market down to a threshold $P_{L,B}^*$ defined where its net present value of leading equals its option value of following, $L_B(P) = F_B(P)$. Firm A, the Leader, possessing the cost advantage g_L , only needs to preempt by an infinitesimally small margin. Firm A enters at $P_{L,B}^*$ but retains its cost advantage:

$$L_A(P_{L,B}^*) = F_B(P_{L,B}^*) + I \cdot g_L \quad (11)$$

This generates the double dividend. The greenium g_L shifts the Leader's entry threshold lower than the standard market threshold, accelerating the initial clean energy transition. Simultaneously, the heavy penalty C_p delays the Follower, significantly extending the Leader's monopoly phase. This combination ensures the Leader's expected value surges, making the bond strictly dominant and proving that decentralized contracting can resolve the coordination failure.

5 Spatial Arbitrage and Extraterrestrial Artificial Intelligence Centers

The baseline model assumes that extraterrestrial energy generation is constrained strictly to near-Earth domains, specifically geostationary orbit. However, the astropolitical frontier offers a spatially differentiated alternative: deep space deployment, specifically at Earth-Moon Lagrange points. By introducing a mutually exclusive spatial choice, the model transforms into a problem of spatial arbitrage under uncertainty.

Deploying to deep space completely resolves the physical tragedy of the commons, as the infinite volume of outer space reduces the endogenous congestion risk to zero ($\lambda_{L1} = 0$). However, the firm faces two severe financial frictions. First, the capital expenditure required to escape the gravity well of Earth is significantly higher. Second, beaming raw power back to Earth over vast interplanetary distances via coherent lasers results in severe transmission

loss, yielding an efficiency discount parameter $\tau \in (0, 1)$.

To circumvent this transmission loss, the firm possesses an embedded real option: rather than beaming power down, the firm can launch the artificial intelligence computation centers directly to the deep space solar array. In this architecture, the solar array powers the supercomputers locally, and the data center utilizes the ambient vacuum of space for costless thermal cooling. The product transmitted back to Earth is not raw energy, but processed computational data, which can be transmitted via laser optical links with negligible energy loss. I model the demand for this space based computation as an independent geometric Brownian motion D_t , evolving with drift α_D and volatility σ_D :

$$dD_t = \alpha_D D_t dt + \sigma_D D_t dW_t^D \quad (12)$$

The value of a firm deploying an integrated artificial intelligence center to deep space is therefore decoupled from the standard terrestrial energy market P_t and relies on a combined cash flow. Because deep space faces zero congestion risk, there is no discrete jump to default risk. The expected present value is strictly the continuous time integral of the two independent stochastic cash flows discounted at the standard risk-free rate r :

$$V_{L1}(P_t, D_t) = \mathbb{E}_0 \left[\int_0^\infty e^{-rt} (\tau \cdot A_{mono}^{osb} P_t + D_t) dt \right] \quad (13)$$

Applying the standard Gordon growth property for geometric Brownian motions, the expectation operators evaluate to $\mathbb{E}[P_t] = P_0 e^{\alpha t}$ and $\mathbb{E}[D_t] = D_0 e^{\alpha_D t}$. Integrating these continuously yields the closed form asset value:

$$V_{L1}(P_t, D_t) = \frac{\tau \cdot A_{mono}^{osb} P_t}{r - \alpha} + \frac{D_t}{r - \alpha_D} \quad (14)$$

Intuitively, this equation proves that while the firm loses efficiency on terrestrial power beaming, it gains a perpetual, risk free exposure to the hyper growth artificial intelligence computation economy.

The mechanism designer can calibrate the orbital sustainability bond not merely to delay the follower, but to induce a spatial pivot. A firm facing the follower role evaluates two mutually exclusive net present values: entering near Earth orbit subject to the penalty C_p , or entering deep space subject to the higher capital cost I_{L1} . The spatial pivot occurs at the terrestrial threshold P_{pivot} where the two values perfectly equalize:

$$V_{GEO}^{osb}(P_{pivot}, \lambda_2) - I_{GEO} - \mathbb{E}[PV(C_p)] = V_{L1}(P_{pivot}, D_0) - I_{L1} \quad (15)$$

By elevating the near Earth congestion penalty C_p to a critical threshold, the expected cost of terrestrial compliance strictly exceeds the transmission and capital frictions of deep space. The optimal strategy of the follower shifts from crowding near Earth orbit to establishing autonomous data centers in deep space. This mathematically achieves a Pareto optimal allocation: it safeguards terrestrial energy safety, utilizes the vacuum of space for cooling, and averts further anthropogenic warming of the Earth biosphere. The secondary firm is incentivized to lay the computational foundation for future space cities, preventing any degradation of the near-Earth commons.

6 Endogenous Risk Mitigation: Reusability and Debris Capture

The standard jump-diffusion literature treats failure rates as exogenous constants. In reality, the advent of reusable rockets, autonomous debris-catching mechanisms, and in-orbit servicing allows firms to actively invest capital to suppress environmental risk.

Let $x \geq 0$ denote the discretionary capital allocated to reusability and debris mitigation. The total sunk cost becomes $I(x) = I_0 + x$. The endogenous congestion risk is a convex, decreasing function of this investment, modeled as $\lambda(x) = \bar{\lambda}e^{-\kappa x}$, where $\kappa > 0$ represents the technological efficacy of robotic systems.

At the moment of entry, the firm must solve a static optimization problem to determine the optimal level of risk mitigation x^* . The firm minimizes its total expected friction, defined as the sum of the upfront mitigation cost x and the expected present value of the congestion penalty. Because the collapse arrival follows a Poisson process with rate $\lambda(x)$, the expected present value of the penalty paid at the random time of collapse τ is $\int_0^\infty e^{-r_L t} \lambda(x) e^{-\lambda(x)t} C_p dt$. Evaluating this integral yields the minimization function:

$$\min_{x \geq 0} \left[x + \frac{\lambda(x)}{r_L + \lambda(x)} C_p \right] = \min_{x \geq 0} \left[x + \frac{\bar{\lambda}e^{-\kappa x}}{r_L + \bar{\lambda}e^{-\kappa x}} C_p \right] \quad (16)$$

To find the optimal investment level x^* , I take the first derivative of the friction function with respect to x and set it to zero. Applying the quotient rule $\left(\frac{u}{v}\right)' = \frac{u'v - uv'}{v^2}$ to the penalty fraction, where $u = \bar{\lambda}e^{-\kappa x}$ and $v = r_L + \bar{\lambda}e^{-\kappa x}$, yields $u' = -\kappa\bar{\lambda}e^{-\kappa x}$ and $v' = -\kappa\bar{\lambda}e^{-\kappa x}$. The derivative of the fraction becomes:

$$\frac{(-\kappa\bar{\lambda}e^{-\kappa x})(r_L + \bar{\lambda}e^{-\kappa x}) - (\bar{\lambda}e^{-\kappa x})(-\kappa\bar{\lambda}e^{-\kappa x})}{(r_L + \bar{\lambda}e^{-\kappa x})^2} C_p = \frac{-r_L \kappa \bar{\lambda} e^{-\kappa x}}{(r_L + \bar{\lambda}e^{-\kappa x})^2} C_p \quad (17)$$

Adding the derivative of the upfront cost (1) yields the final first-order condition:

$$1 - \frac{\kappa \bar{\lambda} e^{-\kappa x} r_L}{(r_L + \bar{\lambda} e^{-\kappa x})^2} C_p = 0 \quad (18)$$

This derivation proves a vital contracting insight: absent the orbital-sustainability bond ($C_p = 0$), the first-order condition collapses to $1 - 0 = 1$. Because the derivative is strictly positive for all x , the cost function is strictly increasing. Therefore, the optimal private investment in reusability and debris capture is cornered precisely at zero ($x^* = 0$). Standard market competition inherently disincentivizes environmental protection because the firm bears the full capital expenditure x but receives zero financial benefit for reducing the global parameter λ .

However, under the bond framework, $C_p > 0$ introduces a heavy financial consequence. The firm endogenously selects an interior optimum $x^* > 0$, perfectly internalizing the externality. The optimal entry threshold P^* is subsequently derived by substituting the minimized cost $I(x^*)$ and optimized risk $\lambda(x^*)$ back into the continuous-time value-matching and smooth-pasting conditions.

7 Venture Capital Contracting and Optimal Capital Structure

The baseline mechanism assumes firms can finance the immense capital expenditure of space infrastructure utilizing corporate balance sheets or green bonds. However, extraterrestrial energy and space based artificial intelligence computation represent frontier technologies with substantial probability of total capital loss. Institutional bondholders strictly avoid this risk profile. Consequently, the pioneer firm must finance the infrastructure through equity issuance, specifically targeting venture capital or sovereign wealth funds.

Let $\eta \in (0, 1)$ represent the fraction of the project equity issued to external investors to raise the required initial capital I . Venture capital funds demand a higher required rate of return r_{vc} to compensate for illiquidity and execution risk, where $r_{vc} > r_L$. The value of the equity claim held by the external investors evaluated at the energy price P_t is the expected present value of their fractional cash flows discounted at their specific hurdle rate:

$$V_{vc}(P_t) = \eta \frac{A_{mono}^{osb} P_t}{r_{vc} - \alpha} \quad (19)$$

To successfully close the funding round, the firm must satisfy the participation constraint

of the investors. The value of the equity claim must equal or exceed the required capital expenditure I . Rearranging this constraint yields the required endogenous dilution fraction $\eta^*(P_t)$ that the founders must surrender as a function of the prevailing market price:

$$\eta^*(P_t) = \frac{I(r_{vc} - \alpha)}{A_{mono}^{osb} P_t} \quad (20)$$

This equation exposes a fundamental financing friction. If the firm attempts to launch early when the energy price P_t is low, the required dilution η^* mathematically exceeds unity, meaning the project is entirely unfinanceable. The founders seek to maximize the absolute value of their retained equity, defined as $W(P_t) = (1 - \eta^*(P_t)) V(P_t)$. Substituting the functional forms of the dilution and the total project value into this objective yields a precise analytical expression for the founder wealth:

$$W(P_t) = \left(1 - \frac{I(r_{vc} - \alpha)}{A_{mono}^{osb} P_t}\right) \frac{A_{mono}^{osb} P_t}{r_L - \alpha} = \frac{A_{mono}^{osb} P_t}{r_L - \alpha} - I \left(\frac{r_{vc} - \alpha}{r_L - \alpha}\right) \quad (21)$$

Because the value of the retained equity $W(P_t)$ is strictly linear with respect to the state variable P_t , the perpetual option to launch the firm held by the founders, $F_W(P_t) = cP_t^\beta$, must satisfy standard intertemporal boundary conditions. At the optimal equity entry threshold P_E^* , the value matching condition $c(P_E^*)^\beta = W(P_E^*)$ and the smooth pasting condition $c\beta(P_E^*)^{\beta-1} = W'(P_E^*)$ must hold simultaneously. Substituting the derivative $W'(P_t) = \frac{A_{mono}^{osb}}{r_L - \alpha}$ into the smooth pasting condition allows the isolation of the constant c . Substituting this constant back into the value matching condition yields the closed form equity entry threshold:

$$P_E^* = \frac{\beta}{\beta - 1} \frac{I(r_{vc} - \alpha)}{A_{mono}^{osb}} \quad (22)$$

This analytical derivation proves that under strict equity financing, the threshold to deploy space infrastructure scales linearly with the venture capital premium r_{vc} . To alleviate this friction and accelerate the deployment of space based artificial intelligence centers, the mechanism designer can introduce a sovereign co-investment contract. If the firm commits to the orbital sustainability bond, a sovereign wealth entity agrees to act as a limited partner in the venture round, mathematically depressing the required fund return r_{vc} . This reduction in the cost of equity narrows the dilution fraction, allowing the firm to initiate the clean energy transition significantly earlier while retaining a controlling interest in the extraterrestrial infrastructure.

8 Two Factor Real Options: Stochastic Demand and Endogenous Launch Costs

The preceding sections utilize a one factor model where the capital expenditure I is deterministic and constant. In reality, the advent of reusable rocketry implies that launch costs are strictly decreasing over time subject to technological volatility. The decision to deploy an artificial intelligence center to space depends not merely on rising terrestrial energy demand, but also on the stochastic decline in payload costs. I extend the framework to a two factor real options model. Let the launch cost I_t follow a geometric Brownian motion with a negative drift μ_I and volatility σ_I :

$$dI_t = -\mu_I I_t dt + \sigma_I I_t dW_I \quad (23)$$

where the Wiener processes for energy demand and launch costs are correlated such that $\mathbb{E}[dW_P dW_I] = \rho dt$.

To determine the optimal value of the unexercised option $V(P_t, I_t)$, the firm must satisfy the two dimensional Hamilton-Jacobi-Bellman partial differential equation. Expanding the expectation operator utilizing the multi-dimensional Ito's Lemma yields:

$$\frac{1}{2}\sigma_P^2 P^2 V_{PP} + \frac{1}{2}\sigma_I^2 I^2 V_{II} + \rho\sigma_P\sigma_I P I V_{PI} + \alpha P V_P - \mu_I I V_I - rV = 0 \quad (24)$$

Because both the payoff and the strike price of the real option are homogenous of degree one in the state variables, the firm evaluates the ratio of the state variables. Let $Z_t = P_t/I_t$ represent the terms of trade between energy revenue and launch costs. Substituting the transformation $V(P, I) = I \cdot v(z)$ allows the extraction of the partial derivatives: $V_P = v'(z)$, $V_I = v(z) - z v'(z)$, $V_{PP} = \frac{1}{I} v''(z)$, $V_{II} = \frac{z^2}{I} v''(z)$, and $V_{PI} = -\frac{z}{I} v''(z)$. Substituting these derivatives back into the two dimensional partial differential equation and dividing the entire equation by I strictly reduces the model to a one dimensional ordinary differential equation:

$$\frac{1}{2}(\sigma_P^2 + \sigma_I^2 - 2\rho\sigma_P\sigma_I) z^2 v''(z) + (\alpha + \mu_I) z v'(z) - (r + \mu_I) v(z) = 0 \quad (25)$$

This rigorous mathematical reduction proves that the drift of the structural ratio is strictly positive, driven simultaneously by rising energy demand α and falling rocket costs μ_I . The firm faces a new characteristic root β_Z derived from the combined volatility profile. Applying the standard value matching and smooth pasting boundary conditions to the ratio variable Z_t , the optimal investment rule is no longer a fixed price scalar, but rather a linear

free boundary in the coordinate space. The firm invests the exact moment the energy price crosses the threshold boundary $P_t = z^* I_t$. This structural proof demonstrates that heavy engineering investments in rocket reusability, which increase the decay parameter μ_I , actively rotate the threshold boundary downward, drastically accelerating the economic viability of space based artificial intelligence data centers.

9 Empirical Calibration Strategy and Methodology

To bridge the gap between continuous-time theory and structural policy implementation, the mathematical parameters of this model must be rigorously calibrated using observable macroeconomic, financial, and aerospace datasets. This calibration strategy anchors the abstract stochastic jump-diffusion processes directly to real-world capital and astropolitical environments.

9.1 Estimating the Geometric Brownian Motion via MLE

The fundamental stochastic variable driving the investment timing game is the terrestrial demand for energy, modeled through the wholesale price P_t . The parameters governing this diffusion—specifically the drift α and the volatility σ —are calibrated using high-frequency wholesale electricity pricing data from the United States Energy Information Administration (EIA), mapping specifically to nodal pricing from the PJM Western Hub.

A primary empirical challenge in electricity markets is that standard hourly or daily spot prices are highly volatile and heavily dominated by transient, mean-reverting weather shocks (e.g., unexpected heatwaves). However, multi-billion dollar extraterrestrial infrastructure investments are not triggered by transient weather noise; they are evaluated against long-term, secular macroeconomic growth. Therefore, to correctly parameterize the continuous-time option value, I apply a 12-month trailing moving average filter to the raw EIA time series. This acts as a mathematical low-pass filter, stripping out short-term mean-reversion and perfectly isolating the secular Geometric Brownian Motion (GBM) expansion trend.

Given that the smoothed underlying price P_t evolves according to a GBM, the discrete log-returns over a monthly time step Δt , defined as $x_t = \ln(P_t/P_{t-\Delta t})$, are independent and normally distributed. Applying standard Maximum Likelihood Estimation (MLE), the discrete sample mean $\hat{\mu}$ and sample variance $\hat{s}^2$ of the log-returns are converted into continuous annualized parameters using Ito’s Lemma:

$$\hat{\sigma} = \frac{\hat{s}}{\sqrt{\Delta t}}, \quad \hat{\alpha} = \frac{\hat{\mu}}{\Delta t} + \frac{1}{2}\hat{\sigma}^2 \quad (26)$$

Executing this estimation on the deseasonalized PJM dataset yields an annualized volatility of $\sigma = 0.1443$ and an annualized drift of $\alpha = 0.0286$. A continuous structural drift of 2.86% accurately encapsulates the sustained demand shock imposed by data center load expansion and AI electrification.

9.2 Parameterizing Financial Constraints and Revenue Premiums

The financial parameters dictating the firm’s intertemporal discount factor are calibrated utilizing macroeconomic data from the Federal Reserve Economic Data (FRED) database. The standard baseline risk-free rate is set to $r = 0.0581$, reflecting the elevated, long-term 10-year Treasury yields required for extended, multi-decade capital horizons. The sustainable finance mechanism introduces an Asymmetric Greenium g , representing the yield concession offered by ESG-mandated institutional investors for sustainable asset deployment. Based on the empirical yield spreads of frontier climate-tech subsidies and green municipal bonds, the greenium is parameterized at 50 basis points ($g = 0.0050$). This effectively reduces the cost of capital for the pioneer firm to an adjusted rate of $r_L = 0.0531$. Total baseline capital expenditure I is calibrated to \$14 Billion ($I = 100$ in normalized model units), utilizing cost-per-kilogram payload metrics published by commercial heavy-lift launch providers such as SpaceX to place a 2-million-kilogram SBSP constellation into operation.

Crucially, the cash-flow multiplier A must reflect the unique physical attributes of space-based energy. Terrestrial solar power is intermittent, suffering from the "duck curve" and yielding lower overall wholesale capture rates. Conversely, a Space-Based Solar Power constellation in geostationary orbit generates uninterrupted, 24/7 carbon-free baseload power. To empirically calibrate this advantage, I utilize Q4 2025 Power Purchase Agreement (PPA) pricing data from LevelTen Energy. Assuming a baseline wholesale average of \$50/MWh, terrestrial solar PPAs command \$61.7/MWh, while 24/7 baseload renewable assets (proxied by advanced wind/geothermal) command \$73.7/MWh.

Therefore, a monopolist SBSP pioneer capturing the 24/7 baseload premium is assigned a revenue multiplier of $A_{mono}^{osb} = \frac{73.7}{50} = 1.474$. If a second firm enters the same orbital sector and causes physical grid congestion, the premium is mathematically collapsed back to the standard intermittent solar proxy, yielding a duopoly multiplier of $A_{duo}^{osb} = \frac{61.7}{50} = 1.234$. This introduces intense endogenous price cannibalization into the model.

9.3 Estimating Endogenous Spatial Risk

Finally, the Poisson arrival rates of an orbital congestion collapse, λ_1 and λ_2 , parameterize the physical astropolitical tragedy of the commons. These structural jump probabilities

are grounded using space situational awareness data and conjunction assessments from the European Space Agency (ESA) Space Debris Office. Let $y_{t,i}$ represent the count of critical orbital conjunctions (close approaches within a 1-kilometer radius) observed during spatial regime i . Assuming these conjunctions follow a non-homogeneous Poisson process, the event rate is scaled by $\phi = 10^{-3}$, a standard probabilistic risk assessment multiplier converting conjunction warnings into terminal Kessler Syndrome probabilities. This empirical framework establishes a baseline monopoly navigation risk of $\lambda_1 = 0.0050$. Upon the entry of a secondary firm, the inability to perfectly coordinate independent, proprietary navigation algorithms triggers an exponential spike in spatial density warnings, empirically shifting the duopoly collapse risk to $\lambda_2 = 0.0600$.

9.4 Calibrating Deep Space Computation and Reusability Parameters

To operationalize the spatial arbitrage and endogenous risk extensions, secondary parameters are calibrated using aerospace and energy agency projections. The growth rate of space based artificial intelligence computation α_D is parameterized at 0.1200, derived from the baseline compound annual growth rate of global data center electricity consumption published in the 2024 International Energy Agency electricity market report. The transmission efficiency of beaming power from deep space τ is set to 0.40, reflecting the inverse square law losses inherent in current microwave array demonstrations.

The capital expenditure penalty for deep space deployment is parameterized utilizing commercial lunar payload services contract data. Because escaping the gravity well requires substantially more delta-v propulsion, the payload cost per kilogram to Lagrange points is empirically observed to be multiple times higher than low Earth orbit. Consequently, I_{L1} is scaled to 2.5 times the baseline near Earth capital expenditure. Finally, the technological efficacy of in orbit servicing and debris capture κ is set to 0.05, proxying the marginal percentage reduction in structural failure probabilities per million dollars invested in autonomous navigation and robotic retrieval systems.

10 Discussion of Numerical Results

Plugging the empirically calibrated vectors into the continuous-time real options mathematical engine yields a precise, dollar-value resolution to the mechanism design problem. By transitioning from abstract algebraic proofs to structural numerical estimation, the profound economic impact of the orbital-sustainability bond becomes distinctly quantifiable.

The outputs are presented comprehensively across three tables and three extensive graphical frameworks, each detailing a distinct structural consequence of asymmetric finance on the astropolitical commons.

10.1 Calibration and Equilibria Shifts

Table 1 explicitly logs the baseline input parameters derived through the empirical calibration strategy. The most critical integration in this matrix is the differentiation between the baseload PPA premium ($1.474x$) and the standard intermittent PPA premium ($1.234x$). By anchoring the continuous-time revenue functions to LevelTen Energy’s empirical spot pricing, the model transitions from a theoretical abstract into a highly actionable corporate finance valuation. Furthermore, identifying the baseline capital expenditure at \$14.00 Billion establishes the massive scale of the sunk costs (I) required to initiate extraterrestrial extraction. At this scale of irreversible capital deployment, even fractional deviations in optimal entry timing result in the loss of hundreds of millions of dollars, necessitating extreme precision in the formulation of the threshold equilibria.

Table 1: Empirical Parameter Calibration (LevelTen Q4 2025 Integration)

Parameter	Description	Value
α	Energy Demand Drift (EIA)	0.0286
σ	Energy Volatility (EIA)	0.1443
r_L	Greenium-Adjusted Rate (50 bps spread)	0.0531
A_{mono}^{osb}	Baseload PPA Premium (Wind Proxy $\frac{73.7}{50}$)	1.474x
A_{duo}^{osb}	Standard PPA Premium (Solar Proxy $\frac{61.7}{50}$)	1.234x
I	Baseline CapEx (SpaceX)	\$14.00B

Note: This table presents the empirically calibrated inputs. Demand parameters are derived via maximum likelihood estimation on PJM wholesale nodal pricing. Revenue multipliers reflect real-world Q4 2025 power purchase agreement spreads, capturing the exact baseload advantage of space-based solar power over terrestrial intermittent renewables.

Table 2 provides the comparative statics of the investment game, proving the fundamental mathematical existence of the pre-emptive race. In a theoretical vacuum where a firm holds a legally protected, unthreatened monopoly right, the optimal macroeconomic entry point is calculated at an energy price of $P_M^* = 8.37/\text{MWh}$. At this specific threshold, the marginal value of waiting perfectly balances the opportunity cost of foregone cash flows. However, because orbit represents an anarchic frontier entirely devoid of enforceable property rights, standard market competition forces the pioneer to pre-emptively deploy capital at a severely depressed price threshold of $P_L^* = 5.33/\text{MWh}$. The firm effectively cannibalizes its own upside to prevent a rival from seizing the market share.

By introducing the optimally calibrated bond, the model successfully engineers a profound double dividend. Subsidized by the 50-basis-point greenium ($r_L = 0.0531$) and the baseload revenue premium, the required hurdle rate of the pioneer drops aggressively from 5.33 down to $P_L^* = 2.98/\text{MWh}$. This achieves the first policy objective: accelerating the immediate deployment of clean energy infrastructure. Simultaneously, the heavy congestion penalty (C_p) inflates the required hurdle rate of the follower from 16.78 up to 21.81. Because the terrestrial energy drift (α) is constrained to a 2.86% continuous annual growth rate, this massive price gap extends the expected deterrence window—the expected timeframe a monopoly remains completely unchallenged—from a standard 40.1 years to a highly secure 69.6 years.

Table 2: Threshold Comparisons: Market Timing under Asymmetric Finance

Metric	Monopoly (Unthreatened)	Standard Race	OSB Optimal Mechanism
Leader Entry (P_L^*)	8.37	5.33	2.98
Follower Entry (P_F^*)	N/A	16.78	21.81
Deterrence Window (Years)	∞	40.1	69.6

Note: The deterrence window calculates the expected hitting time between the entry of the leader and the threshold of the follower based on the geometric drift $\alpha = 0.0286$. The mechanism successfully forces the follower out by nearly three decades, neutralizing the duopoly collision risk.

Table 3 quantifies the exact monetary value of recovering this economic inefficiency. In the standard, unconstrained race, the destruction of option value through premature entry drops the private ex-ante net present value of the firm to \$1.28 Billion. Under the bond market structure, the strategic delay of the follower firmly halts the pre-emptive race, allowing the leader to retain substantial, protected monopoly rents. This dynamic causes the private net present value to surge to \$1.46 Billion. Because the firm earns \$180 Million more under the restrictive environmental covenant than in the free market, adoption of the bond is mathematically proven to be a strictly dominant, subgame perfect strategy. Most importantly, the mechanism does not merely inflate private wealth; the expected present value of the contractual congestion liabilities reserves strictly \$0.41 Billion for the global equity fund. When combining the private NPV gain with the public equity trust, the total expected welfare expands from \$1.28 Billion to \$1.87 Billion, achieving an immense 21.7% recovery ratio of the total potential deadweight loss.

Table 3: Welfare Decomposition and Deadweight Loss Recovery

Metric (Billions USD)	Standard Market	OSB Market	Net Change
Private Firm NPV	\$1.28B	\$1.46B	+\$0.18B
Global Equity Fund	\$0.00B	\$0.41B	+\$0.41B
Total Welfare	\$1.28B	\$1.87B	+21.7% Recovery Ratio

Note: Ex-ante expected values are evaluated at an initial scarcity price of $P_0 = 3.0$. The recovery ratio quantifies exactly how much of the theoretical deadweight loss (the value destroyed by moving from a pure, unthreatened monopoly to a competitive race) is successfully clawed back by the mechanism.

10.2 Visualizing the Strategic Base Mechanism

Figure 1 visually corroborates the structural mechanics of the model across six highly interconnected dimensions. By translating the stochastic calculus into continuous state-space geometries, these panels definitively prove how asymmetric financial contracting governs competitive behavior in anarchic environments.

Panel A illustrates the absolute ex-ante welfare recovery evaluated at the baseline starting price of $P_0 = 3.0$. The solid red bar represents the heavily cannibalized firm value inherent in the standard pre-emptive race. The stacked blue and green bars of the optimal framework graphically prove the mechanism’s success in achieving practical Kaldor-Hicks efficiency. The visualization confirms that not only does the private firm gain net value, but a significant secondary surplus is generated purely for terrestrial social equity. This proves definitively that decentralized, market-driven financial instruments can perfectly substitute for sovereign taxation.

Panel B visualizes the exact mathematical proof of asymmetric rent retention. In a standard unregulated game, the value curve of the leader is forced intertemporally downward until it intersects the option curve of the follower, dissipating all first-mover advantage. However, because the bond grants a 50 basis point greenium and a baseload revenue multiplier exclusively to the certified pioneer, the value curve of the leader is shifted fundamentally higher. Consequently, the leader executes its option exactly where the option of the disadvantaged follower binds at $P = 2.98$. The significant vertical distance between the curves precisely at this threshold represents pure, protected economic rent that the pre-emptive race can no longer dissipate.




Figure 1: Structural Mechanics of the Orbital-Sustainability Bond across Six Dimensions.

Note: Panel A displays Kaldor-Hicks surplus recovery. Panel B provides the mathematical proof of halted pre-emption via asymmetric curves. Panel C demonstrates structural sensitivity to penalty calibration. Panel D simulates the stochastic timeline over a 40-year asset lifecycle. Panel E maps the Inverse Gaussian First Passage Time probability density. Panel F provides the mechanism designer's optimization heatmap for the competitive moat.

Panel C isolates the mathematical sensitivity of the Pigouvian entry delay mechanism. The x-axis tracks the structural magnitude of the congestion penalty (C_p), while the y-axis maps the resultant follower threshold (P_F^*). The strictly monotonic, linear upward slope proves the efficacy of the penalty lever. As the mechanism designer increases the contractual penalty for orbital degradation, the secondary firm is forced to wait for increasingly higher global energy prices to mathematically justify the risk of market entry. This structural geometry allows policymakers to precisely calibrate the penalty required to achieve any desired spatial deterrence window.

Panel D runs a rigorous 40-year Monte Carlo simulation of the energy transition timeline utilizing the empirically calibrated geometric Brownian motion parameters ($\alpha = 0.0286, \sigma = 0.1443$). Starting at $P_0 = 3.0$, the stochastic energy price (the jagged gray line) drifts upward subject to historical volatility. The leader establishes the space-based grid immediately at the blue dashed line (2.98). Over the next 40 years, despite sustained secular macroeconomic growth and highly volatile price peaks, the simulated state trajectory never touches the massive red deterrence barrier (21.81). The follower is actively, physically deterred from entering the orbit, ensuring the spatial state remains locked at $N_t = 1$, thus permanently averting the severe duopoly risk spike (λ_2).

Panel E elevates this analysis by mapping the first passage time probability density function. Utilizing the inverse Gaussian distribution inherent to the optimal stopping times of a geometric Brownian motion, this panel mathematically defines the risk horizon of a terminal Kessler Syndrome collapse. In the standard market (the red shaded area), the probability mass of the follower entering and triggering heavy congestion peaks critically around year 40. By implementing the bond (the green shaded area), the probability mass is structurally flattened and aggressively right-shifted. The statistical probability of triggering catastrophic orbital degradation within the standard operational lifespan of modern infrastructure is reduced to near-zero, successfully preserving the orbital commons for future generations.

Finally, Panel F provides a comprehensive heat map utilized by the mechanism designer to optimize the moat ratio (P_F^*/P_L^*). The x-axis scales the congestion penalty, while the y-axis scales the greenium spread. The dark purple regions in the bottom-left represent critically low moats resulting in heavy pre-emption, while the bright yellow regions in the top-right represent massive deterrence moats yielding near-absolute monopolies. This multi-dimensional matrix proves a critical economic reality: relying solely on a greenium (moving strictly up the y-axis) is insufficient to halt the race. A highly efficient, sustainable transition requires the simultaneous application of both carrots and sticks to push the market into the optimal, highly protected top-right quadrant.

10.3 Visualizing Spatial Arbitrage and Endogenous Mitigation

Figure 2 visualizes the structural dynamics of the spatial arbitrage framework and the endogenous reusability option, explicitly detailing how mechanisms alter the physical trajectory of the space economy.




Figure 2: Spatial Arbitrage, Future Cities, and Endogenous Risk Mitigation.

Note: Panel G identifies the spatial pivot threshold where the congestion cost of near-Earth orbit makes deep space deployment economically superior. Panel H simulates the growth option of the nascent extraterrestrial economy, modeling the demand driver (D_t) for future space urbanization. Panel I visualizes the U-shaped optimization of reusability and ISAM investment (x^*), where the firm balances direct CapEx against expected bond penalties. Panel J illustrates how endogenous mitigation allows for a sustainable reduction in the follower's entry barrier relative to static risk models.

Panel G proves the existence of the spatial pivot. The red line represents the net present value of entering the crowded near Earth orbit subject to the heavy penalty, while the purple line represents the net value of entering deep space. Due to the elevated capital expenditure ($I_{L1} = 250$), the deep space value is initially negative. However, precisely at the threshold where terrestrial energy price crosses $P = 36.8$, the heavy terrestrial penalty C_p mathematically forces the near Earth value curve beneath the deep space curve. At

this exact coordinate, the mechanism successfully repels the follower from near Earth orbit, inducing a Pareto-optimal deep space deployment of artificial intelligence data centers.

Panel H maps the foundational demand driver that makes deep space deployment viable: the extraterrestrial city growth option. Assuming a 12.0% continuous drift (α_D) derived from global data center power consumption rates, the model projects the exponential expansion of computational demand required to sustain lunar mining and orbital habitats. The shaded confidence interval represents the stochastic volatility inherent to frontier market expansion. This demonstrates that deep space deployment is not merely a penalty avoidance strategy, but a highly lucrative long-term real option on the future space economy.

Panel I provides the visual proof of the first-order condition for endogenous risk mitigation. The dashed gray line represents the linear, direct capital expenditure (x) invested in reusable rockets and servicing infrastructure. The dashed red line captures the expected present value of the penalty, which decays exponentially as the physical risk λ decreases. The solid dark line is the total expected friction. Because the mechanism sums a linear cost with a convex benefit, it guarantees a strict, U-shaped global minimum. The firm rationally invests exactly $x^* = 34.0$ to optimize its capital structure, proving that private firms will actively clean the orbit if the bond penalty is rigorously calibrated.

Finally, Panel J demonstrates how endogenous mitigation alters the timeline of the astropolitical transition. If Kessler Syndrome risk is assumed to be a static, exogenous constant, the follower is heavily delayed by the bond, requiring an entry threshold of $P_F = 21.81$. However, by utilizing debris capture technology to internalize the externality, the follower dramatically reduces its physical risk profile. This efficiency increases the present value of the firm, mathematically driving the required entry threshold down to $P_F = 12.10$. This proves that sustainable finance does not strictly halt progress; rather, it incentivizes technological innovations that lower the barriers to sustainable entry.

10.4 Visualizing Venture Capital Dilution and Launch Cost Dynamics

Figure 3 demonstrates the structural realities of financing space infrastructure utilizing equity dilution and two factor stochastic real options boundaries.

Panel K maps the endogenous dilution threshold $\eta^*(P_t)$. The red line represents standard venture capital commanding a 15% required rate of return. At low energy prices, the required dilution exceeds 100%, proving that premature market entry is mathematically unfinanceable. By introducing sovereign wealth co-investments that reduce the hurdle rate to 10% (the green line), the mechanism actively suppresses the required dilution, allowing founders




Figure 3: Venture Capital Contracting and Two-Factor Option Boundaries.

Note: Panel K maps the percentage of the firm founders must surrender to venture capital to fund capital expenditures. Panel L measures the retained equity value. Panel M illustrates the two factor real option boundary where firms evaluate both rising energy prices and falling rocket costs. Panel N simulates the stochastic decline of payload delivery costs via geometric Brownian motion.

to launch infrastructure earlier without surrendering controlling interest in the corporation.

Panel L confirms the maximization problem of the founders. The net present value of the retained equity scales linearly with the energy price. The vertical spread between the two lines perfectly captures the surplus value generated for the founders by securing the subsidized sovereign wealth capital structure.

Panel M provides the visual proof of the two factor partial differential equation. Unlike a standard real option which yields a fixed scalar price threshold, the integration of stochastic launch costs (I_t) yields a linear action boundary defined by $P_t = z^* I_t$. The firm occupies the blue region below the boundary, waiting optimally for either energy demand to rise or rocket prices to fall. The moment the coordinate pair crosses into the red region above the line, the firm exercises the option to launch the artificial intelligence center.

Finally, Panel N visualizes the empirical driver of the I_t state variable. Simulating a geometric Brownian motion with a negative continuous drift ($\mu_I = 0.0400$), the payload cost trajectory mirrors the historical cost reductions achieved by reusable booster technology. This panel confirms that engineering innovations on the launchpad act as a fundamental accelerator for the space based energy transition, independently driving the market toward the structural action boundary.

11 Conclusion

The structural transition to extraterrestrial energy generation is threatened by a dynamic, astropolitical market failure. In the absence of a global sovereign, symmetric pre-emptive competition for finite orbital and planetary resources leads inevitably to premature investment, extreme physical congestion, and the ultimate destruction of economic surplus. This paper proves that this tragedy of the commons can be successfully averted through decentralized, structural mechanism design.

By endogenizing the cost of capital and the environmental preferences of corporate technology buyers, I demonstrate that an orbital-sustainability bond serves as a subgame perfect commitment device. Firms strictly prefer to adopt this highly restrictive instrument over standard, unconstrained finance because the premium baseload revenues generated by constrained technology buyers financially overwhelm the expected cost of an orbital congestion penalty.

When rigorously calibrated using empirical wholesale energy demand drift from the EIA, aerospace CapEx metrics, and an observable 50 basis-point frontier greenium from the Federal Reserve, this mechanism yields a powerful double dividend. It successfully accelerates the deployment of essential clean energy infrastructure, depressing the initial entry threshold to \$2.98/MWh. Simultaneously, it successfully prices the endogenous spatial externality, driving the secondary entrant's threshold out to \$21.81/MWh. As proven by inverse Gaussian first passage time analytics, this extreme Pigouvian delay ensures that secondary competitors do not enter during the primary infrastructure lifecycle, extending the deterrence window to nearly 70 years and effectively averting the exponential spike in Kessler Syndrome probabilities.

Beyond static entry delay, the expanded framework reveals two critical dynamic benefits of the mechanism. First, by properly calibrating the congestion penalty, the mechanism induces a spatial pivot. Secondary firms are mathematically incentivized to bypass crowded near-Earth orbits and instead deploy to deep space, acting as the foundational infrastructure for future extraterrestrial urbanization. Second, when firms are granted the agency to invest

in reusable rockets and debris mitigation, the contract strictly incentivizes them to do so. The financial penalty forces the internalization of environmental costs, leading to an optimal private investment that actively reduces the global failure rate and sustainably lowers long-term entry barriers.

Furthermore, by integrating artificial intelligence computation directly into the deep space arrays, the model demonstrates that firms can bypass the physical transmission limits of energy beaming. By processing data locally in the zero cost thermal environment of deep space and beaming only the resultant data back to Earth, human civilization can structurally decouple computational advancement from terrestrial climate degradation.

By expanding the theoretical framework to include structural venture capital contracting and two factor stochastic boundaries, this research provides a comprehensive blueprint for financing frontier technology. The derivations confirm that sovereign co-investments and the compounding benefits of reusable rocketry actively lower the required equity dilution and accelerate the economically optimal timing of space based infrastructure.

This framework fundamentally halts the value-destroying pre-emptive race. It expands total expected welfare from \$1.28 billion to \$1.87 billion, achieving an immense 21.7% recovery of potential deadweight loss. Furthermore, the mechanism contractually transfers \$410 million of this recovered surplus directly into a global equity fund for terrestrial climate adaptation. Ultimately, this research proves that sophisticated financial contracting can seamlessly align the pursuit of private, astropolitical technological expansion with the preservation of the extraterrestrial commons and the equitable distribution of global resources.

References

- Dixit, A. K. and Pindyck, R. S. (1994). *Investment Under Uncertainty*. Princeton University Press.
- Grenadier, S. R. (2002). Option exercise games: An application to the equilibrium investment strategies of firms. *The Review of Financial Studies*, 15(3):691–721.
- Huisman, K. J. (2013). *Technology investment: A game theoretic real options approach*, volume 28. Springer Science & Business Media.
- Nordhaus, W. D. (2014). Estimates of the social cost of carbon: concepts and results from the dice-2013r model and alternative approaches. *Journal of the Association of Environmental and Resource Economists*, 1(1/2):273–312.
- Ostrom, E. (1990). *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press.
- Pástor, Ľ., Stambaugh, R. F., and Taylor, L. A. (2021). Sustainable investing in equilibrium. *Journal of Financial Economics*, 142(2):550–571.
- Weitzman, M. L. (2009). On modeling and interpreting the economics of catastrophic climate change. *The Review of Economics and Statistics*, 91(1):1–19.